



An Interpretable Dictionary Learning Approach for Abnormality Detection and Segmentation in Chest CT

An undergraduate project report submitted to the

Department of Electrical and Information Engineering
Faculty of Engineering
University of Ruhuna
Sri Lanka

in partial fulfillment of the requirements for the

**Degree of the Bachelor of the Science of Engineering
Honours**

by

K.R.C.S. Kekulawala - EG/2021/4609
B.G.P.L. Thejana - EG/2021/4828
B.G.P.U. Thilakshana - EG/2021/4831
M.D.Y.B. Kularathne - EG/2021/4620

.....
Dr. Kaveen Liyanage
(Supervisor)

.....
Ms. Sithara Mahagama
(Co-Supervisor)

Abstract

Automated abnormality detection from 3D chest Computed Tomography requires clinically interpretable methods to support radiological decision-making. Current State-of-the-Art (SotA) deep learning models often rely on post-hoc explainability techniques that provide approximate localizations and lack direct traceability to the models decisions. Furthermore, SotA models rarely support simultaneous multi-class classification and grounded localization, leading to a fragmented approach to diagnosis. This also hinders the ability to benchmark models across multiple dimensions of performance, including classification accuracy, segmentation quality, and explainability. This study aims to address these challenges through two contributions: First, we introduce an inherently interpretable dictionary learning approach in which detected abnormalities can be back-projected to the original voxel space, allowing diagnostic decisions to be traced back to relevant anatomical structures. Second, we develop a unified, open-source evaluation framework to benchmark SotA models across three axes: (i) classification accuracy, (ii) segmentation quality, and (iii) explainability, using standardized metrics and a publically available 3D chest CT dataset. Our method is benchmarked against SotA models using this framework, demonstrating competitive classification and segmentation performance while providing improved interpretability compared to deep learning baselines.

Acknowledgments

We would like to express our sincere gratitude to Dr. Kaveen Liyanage (Supervisor), Mrs. Sithara Mahagama (CoSupervisor), and Dr. Subhodini Indrapala (Consultant Radiologist at Teaching Hospital, Karapitiya), for their invaluable guidance, support, and constructive feedback throughout the course of this project. Their expertise, mentorship, and encouragement have been immensely helpful in shaping the direction and quality of this research. We are grateful to the Department of Electrical and Information Engineering for providing access to high-performance computing resources and research facilities that were essential for conducting this project. The department's commitment to fostering an environment conducive to research and innovation is greatly appreciated. We would like to acknowledge the creators and maintainers of the publicly available datasets used in this research, particularly the CT-RATE and RexGroundingCT datasets and other medical imaging resources that have made this work possible. These datasets have been crucial for developing and validating the proposed methods. We extend our thanks to the developers and maintainers of the open-source frameworks and libraries used in this project, which significantly contributed in the implementation of this research.

Preface

This thesis is submitted.....

Contents

Abstract	ii
Acknowledgements	iii
Preface	iv
Contents	vii
List of Figures	viii
List of Tables	ix
Acronyms: SOTA – State of the Art	x
1 Introduction	xi
1.1 Background	xi
1.2 Problem Statement	xii
1.3 Objectives and Scope	xii
1.3.1 Objectives	xii
1.3.2 Scope	1
2 Literature Review	2
2.1 Clinical Background	2
2.1.1 Computed Tomography (CT)	2
2.1.2 GGO and Pulmonary Nodules	3
2.2 Computer-Aided Diagnosis (CAD) Systems	4
2.2.1 Evolution of CAD Systems	4
2.3 Deep Learning Approaches	4
2.3.1 nnU-Net	5
2.3.2 BioMedParse	6
2.3.3 TotalSegmentator	7
2.3.4 Merlin	8
2.3.5 MedSAM	9
2.3.6 CT-CLIP	10
2.4 Explainable AI (XAI) in Medical Imaging	11
2.4.1 LIME	11
2.4.2 SHAP	11
2.4.3 Grad-CAM	12
2.4.4 Other CAM Methods	13
2.4.5 Limitations of XAI Methods	14

2.5	Sparse Coding	14
2.5.1	Matching Pursuit (MP)	15
2.5.2	Orthogonal Matching Pursuit (OMP)	15
2.5.3	Least Angle Regression (LARS)	16
2.5.4	Fast Iterative Shrinkage-Thresholding Algorithm (FISTA)	16
2.6	Dictionary Learning (DL)	16
2.6.1	K-Means Singular Value Decomposition (K-SVD)	17
2.6.2	Method of Optimized Directions (MOD)	18
2.6.3	Online Dictionary Learning (ODL)	18
2.7	Discriminative Dictionary Learning	18
2.7.1	Label Consistent K-SVD (LC-KSVD)	19
2.7.2	Fisher Discriminative Dictionary Learning (FDDL)	19
2.7.3	Frozen Dictionary Learning	20
2.8	Dictionary Learning in Medical Imaging	20
2.8.1	Signal Denoising	21
2.8.2	Image Reconstruction and Super-Resolution	21
2.8.3	Detection, Classification and Segmentation	21
3	Methodology	26
3.1	Data Acquisition	26
3.1.1	Datasets	27
3.1.2	Target Abnormality Selection	28
3.1.3	CT Volume Loading	28
3.2	Preprocessing	29
3.2.1	Classical Image Processing-Based Lung Segmentation	30
3.2.2	LungMask-Based Lung Segmentation	31
3.2.3	Patch Extraction	31
3.3	Evaluation Metrics	33
3.3.1	Classification	33
3.3.2	Localization and Segmentation	34
3.3.3	Interpretability of Sparse Dictionary Learning	35
4	Conclusion	36
A	Essential English Grammar - A Handout	37
A.1	Simple Present Tense	37
A.2	Perfect Tense	37
A.3	Past Past Perfect Tense	37
A.4	Passive Voice	37
A.5	Technical Writing Fundamentals	37
A.6	Proof Reading	37
B	Graphical Interpretation of Data	38
B.1	Data Analysis	38
B.1.1	Statistical Analysis	38
B.1.2	Empirical Analysis	38
B.2	The X vs Y vs Z	38
	APPENDIX	39

List of Figures

2.1	Backprojection process in CT imaging	2
2.2	nnU-Net configuration method	5
2.3	Example segmentations by nnU-Net	6
2.4	BiomedParse architecture	7
2.5	Example outputs of BiomedParse	7
2.6	Main classes of TotalSegmentator	8
2.7	Architecture of MedSAM-3	10
2.8	CT-CLIP architecture	10
3.1	Overview of the methodology pipeline.	26
3.2	Shared loading pipeline for CT volumes and finding masks, showing the spacing check, resampling branch (trilinear for volumes, nearest-neighbour for masks) and the preprocessing block applied only to volumes.	29
3.3	Patch extraction workflow, showing the normal-scan sampling grid and the lesion bounding-box branch for annotated scans, each filtered by its own background/lesion-fraction threshold before patches are added to the patch matrix.	33

List of Tables

Acronyms

CXR	- Chest X-Ray
CT	- Computed Tomography
CAD	- Computer-Aided Diagnosis
CNN	- Convolutional Neural Network
ViT	- Vision Transformer
XAI	- Explainable Artificial Intelligence
DL	- Dictionary Learning
LDCT	- Low-Dose CT
GGO	- Ground-Glass Opacity
GGN	- Ground-Glass Nodule
SotA	- State of the Art
PACS	- Picture Archiving and Communication Systems
LIME	- Local Interpretable Model-Agnostic Explanations
SHAP	- SHapley Additive exPlanations
DeepLIFT	- Deep Learning Important FeaTures
cam	- Class Activation Mapping
Grad-CAM	- Gradient-weighted CAM

Chapter 1

Introduction

Developing a unified pipeline for multi-class abnormality classification and grounded localization addresses two of the most critical imperatives in clinical AI adoption: diagnostic accuracy and explainability. Specifically, utilizing dense pixel-level segmentation as the mechanism for localization forces the architecture to provide voxel-level visual evidence for its diagnostic claims. This ensures that the model is accountable and its claims are supported by clinical evidence.

1.1 Background

Lung disease remains a leading cause of global morbidity and mortality [1]. Research shows that early detection of pulmonary diseases significantly improves patient outcomes. For this, radiological imaging tools such as Chest X-Ray (CXR) and Computed Tomography (CT), plays an important role. However, the increasing volume of such imaging data increases radiologists' workload, leading to delays and potential errors in diagnosis. This has motivated the development of automated Computer-Aided Diagnosis (CAD) systems, which use computational algorithms and Artificial Intelligence (AI) to assist in tasks such as abnormality detection, classification, and segmentation [2].

Our research mainly focuses on the development of a CAD system using 3 Dimensional (3D) CT data. While CXR is commonly used as the first imaging test due to its low cost and low radiation exposure, studies show that their sensitivity to abnormalities is limited [3]. Comparatively, CT scans produce high resolution cross-sectional slices that helps detect smaller, early-stage, and subtle lung pathologies that are frequently missed by standard CXR.

However, CT analysis is complex and time-consuming, requiring radiologists to review hundreds of slices per scan, which are often inconsistent due to artifacts, anatomical variations and different scanners/protocols. Furthermore, different anatomical structures can share identical grayscale values. Telling them apart requires expert clinical judgment which is often subjective. High sensitivity can lead to detection of clinically insignificant abnormalities, while visual fatigue can lead to missed findings [4].

To address these challenges, recent research and clinical deployments of AI-based CAD systems are fundamentally changing how radiologists approach CT analysis. AI models such as 3D Convolutional Neural Networks (CNNs) and Vision Transformers (ViTs) have shown to reduce interpretation time, and improve lesion detection sensitivity compared to traditional manual methods. Many of these high perform-

ing models are Deep Learning based, which are criticized for their black box nature, where their internal reasoning is difficult to interpret. This lack of interpretability is a significant barrier for clinical adoption, as radiologists require evidence to support a model’s diagnostic claims.

This has motivated the development of explainable AI (XAI) methods such as Gradient-weighted Class Activation Mapping (Grad-CAM) that helps to visualize exactly which parts of an image a model focuses on when making a specific prediction. Such XAI methods are highly valuable for building clinical trust in deep learning models. However, they often lack the fine-grained, pixel-level precision required to identify tiny lesions, sometimes highlighting incorrect or confounding features [5].

Recent advances in sparse representation models such as Sparse Autoencoders and Dictionary Learning (DL) have shown to provide inherent interpretability in medical imaging. This stems from their ability to reconstruct complex, high-dimensional data as a linear combination of ”atoms” or features. This allows the use of efficient algorithms to easily map the output sparse domain back to the original input feature space, providing a linear relationship between the model’s decision and anatomical structures. Our work builds on these advances by applying DL to 3D Chest CT data for multi-class abnormality classification and grounded localization.

1.2 Problem Statement

The clinical utility of Deep Learning based approaches is hindered by a lack of interpretability. While tools like Grad-CAM provide post-hoc explanations to model decisions, these are often insufficient for medical professionals who require a granular understanding of the model’s decision making process. In high stakes environments, a model that produces a diagnosis without transparency in its reasoning is difficult to integrate into clinical workflows, as it lacks the accountability necessary for legal and ethical compliance.

1.3 Objectives and Scope

The primary goal of this study is to address the limitations of existing AI-based approaches for abnormality classification and localization in chest CTs, by developing an inherently interpretable DL framework. This approach aims to enable transparent, anatomically grounded decision making, where each classification outcome can be traced back to specific pathological patterns. This aligns with clinical reasoning and facilitates validation and adoption in radiological workflows. The specific objectives are as follows:

1.3.1 Objectives

- Implement a Dictionary Learning based CAD system using a publicly available annotated 3D chest CT dataset.
- Map dictionary atoms back to anatomical features and visualize their spatial contribution maps.

- Develop a unified evaluation framework for assessing the performance and interpretability of the proposed approach.
- Benchmark model performance against SotA classification and segmentation models on the same dataset.

1.3.2 Scope

The scope of this study focuses on the development and evaluation of a dictionary learning based approach for classification and grounded localization of lung abnormalities using volumetric 3D CT data, with the following limitations:

- This study is limited to 2 thoracic abnormalities: Pulmonary Nodules, and Ground-Glass Opacities, based on expert radiologist recommendation and dataset availability.
- This study will utilize a single publicly available dataset for training and evaluation, which may limit generalizability to other populations and scanner types.
- The proposed approach will be limited only to interpretable methods, and will not explore approaches that provide better performance for less explainability.

Chapter 2

Literature Review

2.1 Clinical Background

Chronic respiratory diseases account for more than 4 million deaths globally per year [1]. Research shows that early detection of pulmonary diseases, particularly lung cancer and Chronic Obstructive Pulmonary Disease (COPD) dramatically lowers mortality rates. Two landmark studies, the National Lung Screening Trial (NLST) and the Netherlands-Leuven Longkanker Screenings Onderzoek (NELSON) trial, have shown that low-dose CT (LDCT) screening reduces lung cancer mortality by 20% to 33% compared to CXR [6, 7].

2.1.1 Computed Tomography (CT)

CTs generate 3D structural information by acquiring 2D radiographic projections of an object from many angles and computationally reconstructing them into a stack of cross-sectional slices via algorithms such as filtered backprojection [8]. Figure 2.1 shows the Backprojection reconstruction process for a single CT slice.




Figure 2.1: Backprojection reconstruction method for a single slice. (a - Original object slice. b - Set of projections collected at different angles. c - Sinogram resulting from many projections. d - Process of backprojecting the sinogram. e - Final backprojected image. f - Equivalent filtered backprojection image).

Compared with other imaging modalities, CT’s central advantage is its non-destructive, quantitative access to internal 3D structure. CT can visualize and quantify internal features throughout an intact volume, enabling longitudinal or repeated imaging of the same specimen over time. Attenuation-contrast CT, the dominant mode in clinical practice, works well for high-atomic-number, strongly attenuating structures such as bone or calcified tissue and, with contrast agents, vascular or soft-tissue structures. However, it provides relatively poor contrast between materials of similar composition and density, a limitation that phase-contrast and other advanced modalities partially address, though these are less widely deployed clinically due to hardware and dose constraints.

Interpreting 3D chest CT specifically is complicated by several factors described in the imaging literature. Reconstructions are inherently affected by noise and artifacts, motion from breathing or cardiac movement, ring artifacts, and the partial volume effect, where a single voxel spans more than one tissue type, producing intermediate intensities that blur boundaries between structures. These effects broaden and overlap the intensity histograms of different tissues, making simple threshold-based segmentation unreliable and often necessitating boundary-based, region-growing, or learning-based segmentation methods instead. These challenges motivate the shift toward automated, learning-based segmentation approaches capable of handling such structural diversity without extensive manual tuning per dataset.

2.1.2 GGO and Pulmonary Nodules

Our study focuses on the detection and segmentation of two common pulmonary abnormalities: Ground-Glass Opacity (GGO) and pulmonary nodules. A nodule is a discrete, roughly rounded area of increased opacity within the lung parenchyma, generally classified as either a solid nodule or a ground-glass nodule (GGN) based on its internal composition. A GGO refers specifically to a hazy area of increased lung attenuation on CT that does not obscure the underlying bronchial or vascular structures passing through it. GGNs are further subdivided into pure GGNs, which contain no solid component, and mixed (subsolid) GGNs, which contain a solid portion alongside the ground-glass components.

These findings are clinically important because GGNs are frequently associated with early lung cancer and its preinvasive precursors. Notably, GGNs are not uniformly benign or malignant risk indicators equivalent to solid nodules. GGNs are, in fact, more likely to be malignant than solid nodules, and the malignant rate of mixed GGNs is higher than that of pure GGNs.

Early detection of these abnormalities is essential because lung cancer prognosis is highly stage-dependent, and most patients are still diagnosed too late for optimal outcomes. Despite improvements in treatment, the overall five-year survival rate for lung cancer remains only about 10-20% in most countries, largely because most patients are diagnosed at a late stage.

Despite their importance, GGOs and pulmonary nodules remain difficult to detect and characterize reliably. Nodules may be small relative to other anatomical structures in a chest CT, and GGO may be low-contrast against surrounding tissue. These challenges make manual interpretation difficult and motivate the shift toward automated, learning-based segmentation approaches.

2.2 Computer-Aided Diagnosis (CAD) Systems

CAD systems are computer based tools that help medical personnel in analysis of imaging studies and help in making better diagnostic decisions. CAD does not intend to replace doctors and instead, acts as a second opinion by identifying suspicious areas in medical images and providing useful information to support diagnosis. The final diagnosis is always made by the doctor[9].

2.2.1 Evolution of CAD Systems

The modern CAD framework was defined in the 1980s when the University of Chicago shifted the objective from automated diagnosis to acting as a second reader [9]. Early iterations relied on manual image processing using edge detection, filtering, and feature extraction to locate nodules in chest X-rays or microcalcifications in mammograms. Chan et al. provided a pivotal study showing that radiologists using CAD identified more clustered microcalcifications than those working alone, despite a high rate of false positives [10].

Over the 1990s and 2000s, CAD spread to areas such as lung cancer detection on chest X-rays and CT scans, detection of vertebral fractures linked to osteoporosis, brain aneurysm detection using MR angiography, and tracking interval changes in bone scans [9]. A prospective study on mammography, covering tens of thousands of patients, showed a steady rise in cancer detection rates using CAD [11]. This included finding small invasive tumors earlier. Sensitivity for microcalcification clusters rose from roughly 87% in the early 1990s to about 98% by the mid-2000s, while false positives dropped [12]. This "classical CAD" era relied on statistical pattern recognition and manually engineered features. Deep learning has since become the dominant approach in CAD.

The purpose of CAD has also shifted from early tools focused on detection by flagging suspicious zones, to later versions, which added differential diagnosis to estimate if a nodule was malignant or benign. Some systems even retrieve similar past cases from hospital Picture Archiving and Communication Systems (PACS) for comparison. Observer studies indicate that likelihood-of-malignancy scores improve radiologist decisions in subtle cases. Doctors still rely on their own judgment for obvious cases while CAD provides the most utility in borderline situations [?].

Future CAD development will likely move away from stand-alone tools. Instead, these will be integrated packages within PACS, allowing a single scan to be screened for multiple conditions simultaneously [?]. Deep learning has sped up this process. A single trained model can often be adapted for various abnormalities with far less manual effort than previous systems required.

2.3 Deep Learning Approaches

The shift from classical CAD methods to deep learning occurred for a few reasons. Firstly, deep learning models such as CNNs and ViTs have the ability to learn specific image features from pixel data directly rather than having to manually design features and rules for all types of lesions. This makes deep learning models more flexible and often more accurate across different diseases and imaging types. Secondly,

the emerging of large labeled medical imaging datasets and the availability of powerful GPU computing have made it possible to train these deep networks effectively. Thirdly, deep learning networks are more robust to image quality, patient anatomy and scanner types compared to rule based systems. Finally, they continuously outperform classical CAD methods in pivotal studies. The subsequent sections review some of the popular deep learning models used in medical imaging research, and also used in this study.

2.3.1 nnU-Net

The nnU-Net addresses a persistent problem in biomedical image segmentation. Even well-established architectures like the U-Net require extensive, dataset-specific manual tuning of preprocessing, network topology, and training parameters. Small misconfigurations can cause large performance drops. Rather than proposing a new architecture, nnU-Net systematizes this configuration process itself. It decomposes design decisions into three categories:

1. Fixed parameters that generalize across datasets (e.g., architecture template, loss function, optimizer).
2. Rule-based parameters derived automatically from a "dataset fingerprint" (image size, spacing, intensity distribution, modality) via heuristics (e.g., target spacing, patch size, batch size, network depth).
3. A small set of empirical parameters selected via cross-validation (choice among 2D, 3D, and 3D cascade configurations, and whether to apply post-processing such as largest-component suppression).

Figure 2.2 shows the nnU-Net configuration method:




Figure 2.2: nnU-Net configuration method

Evaluated across 23 public datasets and 53 segmentation tasks spanning MRI, CT, electron microscopy, and fluorescence microscopy, nnU-Net outperformed most specialized, task-tuned pipelines without any manual intervention, setting a new state of the art on 33 of 53 target structures. Notably, the authors' analysis of the KiTS

2019 leaderboard showed that configuration choices such as patch size, spacing, normalization affected performance more than architectural modifications. Figure 2.3 shows example segmentations produced by nnU-Net on different datasets:




Figure 2.3: Example segmentations by nnU-Net

This suggests that careful pipeline configuration is often the primary bottleneck in medical image segmentation. This positions nnU-Net less as a competing architecture and more as a strong, reproducible baseline and out-of-the-box framework, and it is frequently cited in the literature as the standard benchmark against which new segmentation methods are compared.

2.3.2 BioMedParse

BiomedParse is a "Foundation Model for Joint Segmentation, Detection, and Recognition of Biomedical Objects Across Nine Modalities" [13]. Rather than relying on a fixed label set during inference, the model accepts a natural language prompt describing the target anatomy or pathology and produces a pixel level mask for the queried concept. This makes BiomedParse especially useful in settings where the structures of interest are heterogeneous, difficult to enumerate in advance, or described differently across institutions and reports.

The model combines image encoding with language understanding and uses a transformer based decoder to align the prompt with the visual features extracted from the scan [13]. Figure 2.4 shows the flowchart of the architecture:




Figure 2.4: BiomedParse architecture

The output is a segmentation mask, as shown in Figure 2.5.




Figure 2.5: Example outputs of BiomedParse

BiomedParse serves as a SotA reference model for discussing localization quality, and interpretability in our study.

2.3.3 TotalSegmentator

Many deep learning methods were developed to segment only one or a few organs. Wasserthal et al. introduced TotalSegmentator, a deep learning framework that can automatically segment 104 anatomical structures from a single CT scan [14]. Figure 2.6 shows the main classes of TotalSegmentator, which include 27 organs, 59 bones, 10 muscles, 8 vessels.




Figure 2.6: Main classes of TotalSegmentator

The framework is built on the nnU-Net architecture and was trained using more than 1,200 routine clinical CT scans collected from different hospitals and patient groups [14]. It achieved a mean Dice Similarity Coefficient (DSC) of about 0.94. Because of its high segmentation accuracy, TotalSegmentator is commonly used as a preprocessing tool in medical image research [15].

2.3.4 Merlin

Merlin is an AI model that can help interpret CT scans automatically. 3D vision-language model (VLM) built to understand abdominal CT scans by learning jointly from the images, radiology reports, and electronic health record (EHR) data [16].

Most earlier medical vision-language models were built for 2D images, such as chest X-rays, paired with short captions. CT scans are different, they are 3D volumes made of hundreds of slices, and simply applying a 2D model slice-by-slice throws away important 3D spatial information, like how a lesion connects across neighboring slices. Merlin was built to fix this by processing the full 3D CT volume directly, instead of treating it as a stack of unrelated 2D images [16].

Merlin is trained using over 6 million CT images from more than 15,000 scans, paired with over 1.8 million EHR diagnosis codes and more than 6 million tokens of radiology report text [16]. Instead of relying only on manually labeled data, which is slow and expensive to collect, Merlin learns from data hospitals already generate: doctors' reports and diagnosis codes. Merlin follows the concept introduced by CLIP, a foundational vision-language model designed originally for natural images. In CLIP style training, similar or matching image or text pairs closer to each other in a shared representation space while non-matching pairs are pushed apart making the model learn which images and text describe the same thing. Merlin adapts this idea to 3D CT data and adds EHR code supervision on top, which earlier medical VLMs generally did not use [16].

Merlin is part of a broader wave of work bringing vision-language modeling into 3D medical imaging, alongside similar efforts like CT-CLIP, which uses the same contrastive approach but focuses on chest CT scans[16]. Overall, Merlin shows that

learning directly from full 3D volumes and everyday hospital data, rather than small 2D datasets with manual labels, can produce useful for reducing radiologist workload.

2.3.5 MedSAM

MedSAM adapts the Segment Anything Model (SAM), originally developed by Meta AI for natural images, into a general-purpose segmentation model for medical scans [17]. Rather than training a separate network for each organ or modality, a single fine-tuned model is reused across many segmentation tasks. A key limitation is that MedSAM operates on 2D slices. Applying it to a 3D chest CT volume therefore requires a separate bounding box on every slice, and the resulting masks do not capture how a nodule connects across neighbouring slices.

MedSAM2 addresses this limitation by building on the SAM2 architecture and its memory-attention mechanism, allowing a single prompt to be propagated across an entire volume instead of being redrawn slice by slice, which substantially reduces the annotation effort required [18]. Specifically, that prompt is purely spatial, a box(two corner points), a point(roughly inside the target), or a mask(rough/full segmentation on that one slice) on one key slice, with no text or class label involved; from it, MedSAM2 segments the object on that slice and then propagates the mask forward and backward through the remaining slices via memory attention, without requiring a new prompt on each one.

Despite this improvement, pulmonary findings remain difficult for these models to segment accurately. A benchmarking study comparing MedSAM and MedSAM2 against standard segmentation networks found that both achieved strong performance on whole-lung segmentation but performed noticeably worse on individual tumours and small nodules, with results sensitive to the quality and consistency of the box prompts [19]. This reflects a broader limitation of prompt-driven segmentation models.

This sensitivity to prompt quality points to a more basic limitation: neither MedSAM nor MedSAM2 can be told what to segment in words, since the underlying backbones have no language pathway. Meta’s SAM3 targets this gap directly through what it calls Promptable Concept Segmentation, allowing a target to be specified by a short noun phrase rather than only by a box, point, or mask [20]. MedSAM-3 takes this a step further and fine-tunes the architecture on medical images paired with concept phrases, such as lung infection for chest X-ray or pulmonary artery for lung CT [21]. Even so, the authors’ own ablations make clear that text alone does not remove the need for a geometric anchor: on the 3D Parse2022 lung CT benchmark, raw SAM 3 given nothing but the phrase ”pulmonary artery” managed a Dice score of only 0.53, well short of nn-U-Net’s 0.83, while on 2D benchmarks such as BUSI, text-only prompting scored 0 for SAM 3 and just 0.27 after MedSAM-3 fine-tuning, jumping to 0.78 only once a bounding box was added alongside the text [21]. It is also worth noting that MedSAM-3’s fine-tuning was limited to four 2D datasets, and its 3D lung CT evaluation went no further than the untuned SAM 3 baseline on pulmonary artery delineation; neither this paper nor the wider MedSAM family reports results on lung nodule or ground-glass-opacity segmentation in chest CT.




Figure 2.7: Architecture of MedSAM-3

2.3.6 CT-CLIP

CT images contain valuable information for diagnosing pulmonary diseases and are commonly used in many Deep Learning approaches as a single source of truth. However, radiology reports also provide detailed clinical descriptions of the findings in these CT images. Hamamci et al. proposed Computed Tomography Contrastive Language-Image Pre-training (CT-CLIP) [22], a vision-language foundation model for chest CT images, which utilizes both modalities. Unlike traditional supervised models that require large amounts of labeled data, CT-CLIP learns the relationship between CT scans and their corresponding radiology reports using contrastive learning. This allows the model to learn meaningful image and text representations without task-specific annotations.

CT-CLIP consists of a 3D image encoder and a text encoder. The image encoder extracts features from CT volumes, while the text encoder processes the corresponding radiology reports. Figure 2.8 shows the architecture of CT-CLIP.




Figure 2.8: CT-CLIP architecture

During training, the model learns to match each CT scan with its correct report in a shared feature space. The model was trained using the CT-RATE dataset [23], which pairs over 50,000 reconstructed 3D chest CT volumes from around 25,700 scans of over 21,000 patients with their corresponding radiology reports. It has shown

strong performance in pulmonary disease classification, image retrieval, and zero-shot learning tasks, outperforming fully supervised models across key metrics [22].

2.4 Explainable AI (XAI) in Medical Imaging

Many deep learning models achieve high accuracy, but their decision-making process is often difficult to interpret and are referred to as "black-box" models [24]. In healthcare, this lack of transparency can reduce clinicians' confidence in AI-assisted diagnosis, as medical decisions directly affect patient outcomes. XAI methods aim to solve this issue.

2.4.1 LIME

Local Interpretable Model-agnostic Explanations (LIME), proposed by Ribeiro et al. [25], is a model-agnostic technique for explaining the predictions of any classifier or regressor by locally approximating its behavior with an interpretable surrogate model. The core idea is to treat the original model f as a black box and, for a given prediction on instance x , learn a simple explanation model g (typically a sparse linear model) that is faithful to f in the vicinity of x , even if g cannot approximate f well globally. Formally, LIME finds an explanation by minimizing an objective:

$$\xi(x) = \arg \min_g L(f, g, \pi_x) + \omega(g)$$

where L measures how unfaithful g is to f within a local neighborhood weighted by a proximity function π_x , and $\Omega(g)$ penalizes the complexity of the explanation to keep it human-interpretable. To ensure explanations are understandable regardless of the model's actual features, LIME operates over an interpretable representation such as a set of super-pixels for images—rather than the model's raw internal features (pixel tensors).

Through simulated experiments and studies, the authors demonstrate that LIME explanations are faithful to the underlying model, recovering 90% of known important features [25]. It helps users select the better-generalizing classifier between two competing models even when held-out accuracy is misleading, enables non-experts to perform effective feature engineering to improve an untrustworthy classifier, and allow users to detect models that rely on spurious correlations.

2.4.2 SHAP

SHapley Additive exPlanations (SHAP) is a unified framework for interpreting the predictions of complex machine learning models, proposed by Lundberg and Lee [26]. It builds on the observation that many existing interpretability methods—including Local Interpretable Model-agnostic Explanations (LIME), Deep Learning Important Features (DeepLIFT), layer-wise relevance propagation, and classical Shapley value approaches—can all be framed as instances of a single class called additive feature attribution methods, in which an explanation model g approximates the original model f as a linear function of binary "presence/absence" indicator variables:

$$g(z') = \phi_0 + \sum \phi_i z'_i$$

Within this class, the authors show that only one explanation satisfies three desirable properties simultaneously:

1. Local Accuracy - The explanation must sum to match the model’s actual output.
2. Missingness - Features absent from the input receive zero attribution.
3. Consistency - A feature’s attributed importance cannot decrease if the model comes to rely on it more, holding other inputs fixed.

This unique solution fairly distribute a prediction’s output among input features by averaging each feature’s marginal contribution across all possible orderings/subsets of the other features—giving rise to the name SHAP values.

Because exact Shapley value computation is combinatorially expensive, the paper introduces several practical estimation methods. Kernel SHAP is a model-agnostic approach that computes SHAP values via specially weighted linear regression, connecting the LIME framework to Shapley values by deriving the specific loss function, weighting kernel, and regularization needed to make LIME’s regression-based estimate satisfy the above three properties [26].

For deep neural networks specifically, the authors introduce Deep SHAP, which combines the compositional, layer-by-layer propagation style of DeepLIFT with Shapley value theory to more efficiently approximate feature attributions, and Max SHAP, which computes Shapley values for max-pooling operations in polynomial rather than exponential time. Empirically, the authors demonstrate that SHAP values align more closely with human-generated explanations than LIME or DeepLIFT, and that Kernel SHAP achieves comparable accuracy to sampling-based Shapley estimation with substantially fewer model evaluations, establishing SHAP as a theoretically grounded and practically efficient tool for model interpretability [26].

2.4.3 Grad-CAM

Gradient Class Activation Mapping (Grad-CAM), proposed by Selvaraju et al. [27], is a technique for producing visual explanations from CNNs by using gradient information to identify which regions of an input image are important for a given prediction. Unlike its predecessor, Class Activation Mapping (CAM), Grad-CAM works with any CNN-based architecture without changes or retraining. The method computes the gradient of a target class score with respect to the feature maps of the last convolutional layer, then global-average-pools these gradients across the spatial dimensions to obtain a set of neuron importance weights α_k^c . A weighted combination of the forward activation maps using these weights are passed through a ReLU to retain only features with a positive influence on the target class. This produces a coarse heatmap highlighting class-discriminative regions of the image. The authors formally prove that Grad-CAM is a strict generalization of CAM, reducing to the same weights when applied to CAM-compatible architectures.

Through extensive evaluation, the authors demonstrate that Grad-CAM achieves superior weakly-supervised localization performance compared to prior methods without sacrificing classification accuracy [27]. They further showed that Grad-CAM’s localizations correlate more strongly with occlusion-based faithfulness measures than competing visualization techniques, indicating it is more faithful to the underlying model. Beyond diagnostic use for uncovering CNN failure modes and robustness to

adversarial perturbations, the authors demonstrate Grad-CAM’s utility for identifying and mitigating dataset bias, famously showing that a biased ”doctor vs. nurse” classifier had learned to rely on gender-correlated cues such as hairstyle, rather than clinically relevant features.

2.4.4 Other CAM Methods

Grad-CAM++

Chattopadhyay et al. propose Grad-CAM++ [28] as a generalization of Grad-CAM that addresses two of its key limitations:

1. Degraded localization when multiple instances of the same class appear in a single image
2. Incomplete coverage of an object’s extent in single-object images

While gradient-based methods like Grad-CAM provide fine-grained explanations, their performance drops when localizing multiple occurrences of the same class, and Grad-CAM heatmaps often fail to capture the entire object [28]. Grad-CAM++ uses a weighted combination of the positive partial derivatives of the last convolutional layer’s feature maps with respect to a specific class score as pixel-wise weights, producing a visual explanation for the class of interest. By deriving closed-form higher-order gradient weighting coefficients rather than Grad-CAM’s simple global-average-pooled gradients, Grad-CAM++ generates explanation maps that better localize multiple object instances and cover objects more completely.

Score-CAM

Score-CAM proposed by Wang et al. [29] removes the dependence on gradients entirely. Unlike previous CAM approaches, Score-CAM obtains the weight of each activation map through its forward-passing score on the target class rather than through gradients, with the final saliency map obtained as a linear combination of these weights and the activation maps. Concretely, each channel of the final convolutional layer’s activation map is upsampled to input resolution and used as a mask over the original image; the resulting masked images are forward-passed through the network, and the change in the target class’s confidence score serves as that channel’s importance weight. Because it dispenses with backpropagated gradients—which can be noisy or saturate—Score-CAM was shown to produce better visual quality and fairness in explanation than gradient-based CAM variants, at the cost of requiring multiple forward passes.

LayerCAM

Jiang et al. introduced LayerCAM [30] to address the coarse spatial resolution of standard CAM/Grad-CAM maps, which are typically generated only from the final convolutional layer. Because such maps have low spatial resolution, they often only locate coarse regions of target objects, limiting performance on tasks needing pixel-accurate localization. By rethinking the relationship between feature maps and their corresponding gradients, LayerCAM produces reliable class activation maps at different layers of the CNN, using element-wise gradient weighting, rather than globally

pooled gradients, so that shallower layers—which retain finer spatial detail—can also contribute meaningful, class-discriminative localization. These layer-wise maps can then be fused into a single higher-resolution activation map, better capturing fine object boundaries than Grad-CAM’s coarse, final-layer-only heatmap.

2.4.5 Limitations of XAI Methods

Despite substantial progress, XAI continues to face a fundamental accuracy-interpretability tradeoff: highly interpretable models such as decision trees, linear regression, and rule-based learners are intrinsically limited in their ability to model non-linear, high-dimensional relationships. Deep neural networks that achieve state-of-the-art performance remain the least interpretable [31]. Model extraction and distillation techniques, which attempt to approximate black-box behavior with a simpler surrogate model, partially address this but reduces fidelity to the original model’s decision boundaries. Post-hoc explanation methods such as LIME offer only local fidelity and provide an unsatisfactory global approximation, and saliency maps such as Grad-CAM can produce noisy uninterpretable outputs. Furthermore, saliency maps assume that all model-relevant features are inherently human-interpretable, which does not always hold [32]. Explanation techniques are also vulnerable to manipulation and adversarial perturbation; LIME has been shown to develop self-learned patterns that produce misleading explanations under such conditions, raising concerns about the robustness and trustworthiness of the interpretations themselves [32].

Beyond technical constraints, XAI suffers from a lack of standardized, universally accepted evaluation frameworks, since interpretability and explainability are frequently conflated and no consensus exists on how to rigorously quantify explanation quality [33, 34]. Existing evaluation paradigms: application-grounded, human-grounded, and functionally grounded, each carry tradeoffs between cost, human involvement, and validity. The subjective, user-dependent nature of what constitutes a “good” explanation makes benchmarking especially difficult [31]. Additional challenges include the computational expense of methods like Model Class Reliance, the difficulty such methods have scaling to complex, high-dimensional, or strongly feature-correlated models, and the broader absence of a unified theoretical framework connecting the disparate strands of explanation research [35].

2.5 Sparse Coding

Why is sparsity useful?

How is sparse representation learned?

Sparse coding is a signal representation framework in which a signal is expressed as a linear combination of only a small number of elements (“atoms”) drawn from a dictionary, rather than as a dense combination over an orthogonal basis. Formally, given a signal x and a dictionary D whose columns are atoms, sparse coding seeks a coefficient vector α such that:

$$x \approx D\alpha,$$

with α containing mostly zero entries. A key aspect of this framework is the use of overcomplete representations, where the dictionary D has more atoms than the

dimensionality of the signal itself (i.e., more columns than rows). This redundancy is deliberate: it gives the representation greater flexibility to match the underlying structure of the signal, allowing different signals to be well-approximated by different small subsets of atoms, which improves representational power, robustness to noise, and the ability to capture diverse structural patterns compared to a complete (non-redundant) basis. The trade-off is that recovering α becomes an underdetermined and generally NP-hard combinatorial problem, since many possible coefficient vectors could reconstruct $\mathbf{x}$, and one must impose sparsity as a constraint or penalty to select a meaningful solution.

Because exact sparsest-solution search is computationally intractable, two main families of practical algorithms are used:

1. Greedy Algorithms (Matching Pursuit, OMP)
2. Convex relaxation Algorithms (LARS, FISTA)

Greedy methods iteratively select the atom most correlated with the current residual, add it to the active support set, and update the coefficients, repeating until a stopping criterion (sparsity level or residual error) is met; these methods are fast and simple but only approximate the optimal solution.

Convex relaxation methods instead replace the intractable ℓ_0 -norm (which counts nonzero entries) with the ℓ_1 -norm, yielding a convex optimization problem—as in Basis Pursuit or the Lasso—that can be solved reliably and, under certain conditions on the dictionary (e.g., restricted isometry or incoherence properties), provably recovers the same sparse solution as the original ℓ_0 problem. In practice, greedy methods are often preferred for speed and low-sparsity regimes, while convex relaxation approaches tend to offer stronger theoretical guarantees and more stable solutions in noisier or higher-sparsity settings.

2.5.1 Matching Pursuit (MP)

Matching Pursuit is a greedy algorithm for approximating a signal $\mathbf{x}$ as a sparse combination of atoms from a dictionary $\mathbf{D}$. It proceeds iteratively: at each step, it computes the inner product between the current residual (initially the signal itself) and every atom in the dictionary, selects the atom with the highest absolute correlation, and updates the coefficient for that atom by projecting the residual onto it. The residual is then updated by subtracting this contribution, and the process repeats until a stopping criterion is reached (target sparsity or residual energy threshold). MP is simple and computationally cheap, but because it only updates one coefficient per iteration without adjusting previously selected coefficients, it can select the same atom multiple times and converges slowly, and its solutions are generally suboptimal compared to more refined methods.

2.5.2 Orthogonal Matching Pursuit (OMP)

OMP improves on MP by re-optimizing all currently selected coefficients at every iteration rather than only the most recently added one. At each step, it selects the atom most correlated with the residual (as in MP), adds it to the active support set, and then solves a least-squares problem over all atoms currently in the support to

obtain new coefficients that minimize the residual jointly. Because the residual is then orthogonal to all previously selected atoms, OMP never reselects the same atom twice, and it typically converges in fewer iterations and yields more accurate reconstructions than plain MP. The added cost is the least-squares solve at each iteration, making OMP more computationally expensive per step, though still tractable for moderate sparsity levels.

2.5.3 Least Angle Regression (LARS)

LARS is an efficient algorithm for computing the entire regularization path of the Lasso (ℓ_1 -penalized least squares) problem. Rather than fixing a penalty parameter λ and solving a single optimization, LARS starts with all coefficients at zero and iteratively identifies the predictor (atom) most correlated with the current residual, moving the coefficient in that predictor's direction. Instead of moving all the way to the least-squares solution as in stepwise regression, LARS moves only until another predictor becomes equally correlated with the residual, at which point that predictor is added to the active set and the direction is updated to remain equiangular between all active predictors. With a simple modification (allowing coefficients to be dropped from the active set when they cross zero), LARS exactly reproduces the full Lasso solution path with a computational cost comparable to ordinary least squares, making it a popular and efficient tool for convex-relaxation-based sparse coding, particularly when the full path of solutions across sparsity levels is of interest.

2.5.4 Fast Iterative Shrinkage-Thresholding Algorithm (FISTA)

FISTA solves ℓ_1 -regularized sparse coding problems, where the goal is to minimize:

$$\sqrt{\|x - D\alpha\|^2} + \lambda\|\alpha\|_1$$

It builds on the Iterative Shrinkage-Thresholding Algorithm (ISTA), which alternates between a gradient descent step on the data-fidelity term and a soft-thresholding (shrinkage) step that enforces sparsity by shrinking small coefficients toward zero. While ISTA converges at a sublinear rate of $O(1/k)$, FISTA accelerates this by incorporating a momentum term—extrapolating using a weighted combination of the current and previous iterates—which improves the convergence rate to $O(1/k^2)$ without increasing per-iteration cost. This makes FISTA well-suited to large-scale sparse coding and compressed sensing problems where fast, reliable convergence to the ℓ_1 -regularized solution is needed.

2.6 Dictionary Learning (DL)

Dictionary learning is a technique in which an image patch or a region of a medical scan is represented as a combination of a small number of basic building blocks, rather than by describing every pixel value directly. This idea comes from sparse coding theory and is widely used in image denoising, compression, and interpretable machine learning, including medical image analysis [36] [37].

A dictionary is a matrix, written as $D \in \mathbb{R}^{n \times k}$, where each column d_i is called an atom. Each atom is a small, fixed pattern, for example a short edge, a texture, or a simple shape, that acts like a basic building block for describing images [36].

$$D = [d_1, d_2, \dots, d_k], \quad \|d_i\|_2 = 1 \quad \forall i$$

Each atom is normalized to unit length so that none of them dominate because of scale. A good dictionary has atoms that are simple and reusable.

Once a dictionary is available, any image $x \in \mathbb{R}^n$ is approximated as a weighted sum of a small number of atoms:

$$x \approx D\alpha = \sum_{i=1}^k \alpha_i d_i$$

Here, $\alpha \in \mathbb{R}^k$ is the sparse code; most of its values are zero, and only a few atoms are actually used to rebuild the image. Finding this code is called sparse coding, and it can be solved by minimizing the reconstruction error together with a sparsity penalty [37] [38]:

$$\min_{\alpha} \frac{1}{2} \|x - D\alpha\|_2^2 + \lambda \|\alpha\|_1$$

This is the Lasso formulation, where λ controls how sparse the solution is [39]. When the dictionary itself is learned from training data $X = [x_1, \dots, x_m]$, both D and the codes $A = [\alpha_1, \dots, \alpha_m]$ are optimized together

$$\min_{D, A} \|X - DA\|_F^2 + \lambda \|A\|_1 \quad \text{subject to} \quad \|d_i\|_2 = 1$$

Learned atoms are not random they turn out to look like meaningful patterns. Since each image patch is reconstructed using only a small number of atoms. Each atoms tends to specialize in representing a simple pattern rather than blending many patterns together. This lets a person actually look at a learned dictionary and understand what each atom represents, which is much harder with internal weights of a deep neural network. In medical imaging, atoms can end up representing specific tissue textures or lesion patterns that a radiologist can visually recognize[40].

Learning a dictionary is usually done by fixing the dictionary and solving for sparse codes, then fixing the codes and updating the dictionary. The sparse coding step is commonly solved with Orthogonal Matching Pursuit (OMP), a greedy method that adds one atom at a time [41]

$$i^* = \arg \max_i |d_i^T r|, \quad r = x - D\alpha$$

or with Lasso/LARS, using the ℓ_1 -penalized objective shown above [39]. The dictionary update step is where the main algorithms differ, covered below.

2.6.1 K-Means Singular Value Decomposition (K-SVD)

K-SVD is one of the most widely used algorithms for learning a dictionary, and it solves the following optimization problem[39]

$$\min_{D, X} \{\|Y - DX\|_F^2\} \quad \text{subject to} \quad \forall i, \|x_i\|_0 \leq T_0$$

Here, Y is the matrix of training data, for example CT image patches, D is the dictionary being learned, and X is the matrix of sparse codes, where each column x_i is the code for one training sample. The constraint $\|x_i\|_0 \leq T_0$ means that each

code is allowed to use at most T_0 atoms, using the ℓ_0 norm to directly count non-zero entries rather than approximating sparsity with an ℓ_1 penalty like Lasso does. This gives K-SVD tighter control over exactly how sparse each representation is.

K-SVD solves this problem by alternating between two steps. First, with the dictionary D fixed, it finds the sparse codes X using a pursuit algorithm such as Orthogonal Matching Pursuit. Second, with the codes fixed, it updates the dictionary one atom at a time instead of all at once, which makes it converge faster and produce sharper, more specific atoms [39].

2.6.2 Method of Optimized Directions (MOD)

Method of Optimized Directions is an earlier and simpler approach than K-SVD. Instead of updating atoms one at a time, MOD updates the entire dictionary at once, using a direct least-squares solution [?]. Given the training data X and current sparse codes A , the optimal dictionary is:

$$D = X A^T (A A^T)^{-1}$$

This comes directly from setting the derivative of the reconstruction error $\|X - D A\|_F^2$ to zero and solving for D . MOD is fast and easy to implement, but because it updates all atoms together, it tends to need more iterations to converge compared to K-SVD, and can get stuck producing less sharp atoms [?].

2.6.3 Online Dictionary Learning (ODL)

Both K-SVD and MOD assume the entire training dataset fits in memory and is processed all at once, which becomes impractical for very large image datasets. Online Dictionary Learning solves this by processing the data in small batches, updating the dictionary incrementally as new samples arrive, instead of waiting to see the whole dataset [?]. At each step t , a new sample x_t is drawn, its sparse code α_t is computed using the current dictionary, and then the dictionary is updated using accumulated statistics from all samples seen so far:

$$A_t = A_{t-1} + \alpha_t \alpha_t^T, \quad B_t = B_{t-1} + x_t \alpha_t^T$$

$$D_t = \arg \min_D \frac{1}{t} \sum_{i=1}^t \left(\frac{1}{2} \|x_i - D \alpha_i\|_2^2 \right)$$

which can be solved efficiently using A_t and B_t instead of storing all past samples. This makes ODL well suited for large-scale problems, including learning dictionaries from thousands of medical image patches, since memory use stays constant regardless of dataset size [?].

2.7 Discriminative Dictionary Learning

Traditional dictionary learning methods mainly focus on minimizing the reconstruction error between the input data and its sparse representation [42]. However, good reconstruction does not necessarily produce features that are useful for classification.

Samples from different classes may activate similar atoms if the learned dictionary only captures common structures in the data.

Discriminative dictionary learning addresses this limitation by incorporating class information during dictionary optimization. The objective is not only to represent the input samples accurately but also to learn sparse codes that improve the separation between different classes. By learning class aware atoms, discriminative dictionary learning can provide both classification capability and a more interpretable representation. Several approaches have been proposed, including Label Consistent K-SVD (LC-KSVD) and Fisher Discriminative Dictionary Learning (FDDL), which extend the K-SVD framework by introducing supervised learning constraints.

2.7.1 Label Consistent K-SVD (LC-KSVD)

Label Consistent K-SVD (LC-KSVD), introduced by Jiang et al. [43], extends the classic K-SVD framework by embedding label consistency into the dictionary update process. The core idea is to encourage sparse codes from samples of the same class to exhibit similar activation patterns, while those from different classes diverge. This is achieved without sacrificing the algorithm's efficient iterative structure involving sparse coding and dictionary refinement.

The method comes in two main variants. LC-KSVD1 incorporates a label consistency term that aligns the sparse codes with an ideal discriminative structure through a linear transformation:

$$\min_{D, A, A_c} \|X - DA\|_F^2 + \alpha \|Q - A_c A\|_F^2 \quad \text{s.t.} \quad \|a_i\|_0 \leq T \quad \forall i. \quad (2.1)$$

Here, Q is a target matrix encoding class specific indicator patterns, and A_c learns to map codes toward this structure. LC-KSVD2 extends LC-KSVD1 by adding a classifier learning term. The dictionary, sparse codes, label consistency transformation, and classifier are jointly optimized:

$$\min_{D, A, A_c, W} \|X - DA\|_F^2 + \alpha \|Q - A_c A\|_F^2 + \beta \|H - WA\|_F^2 \quad \text{s.t.} \quad \|a_i\|_0 \leq T, \quad (2.2)$$

where H contains the label information. The additional classification term allows the learned sparse representation to be directly optimized for prediction. Due to its ability to learn class specific atoms, LC-KSVD has been applied in image classification tasks where interpretability of learned features is important [43].

2.7.2 Fisher Discriminative Dictionary Learning (FDDL)

Fisher Discriminative Dictionary Learning (FDDL), proposed by Yang et al. [44], is another supervised dictionary learning approach that directly enforces Fisher like criteria on both the dictionary atoms and the sparse coefficients. It learns a structured dictionary $D = [D_1, D_2, \dots, D_c]$ with class specific sub dictionaries D_i while promoting sparse codes A that exhibit low within class scatter and high between class scatter.

The core objective combines a discriminative fidelity term $r(\cdot)$ (ensuring good reconstruction within class and poor reconstruction across classes) with a Fisher discrimination penalty on coefficients:

$$\min_{D, X} \sum_i r(A_i, D, X_i) + \lambda_1 \|X\|_1 + \lambda_2 \left(\text{tr}(S_W(X)) - \text{tr}(S_B(X)) + \eta \|X\|_F^2 \right), \quad (2.3)$$

where $r(\cdot)$ encourages each class to be well represented by its own sub dictionary while poorly represented by others. S_W and S_B denote the within class and between class scatter matrices of the sparse codes. The elastic term controlled by η is added to guarantee that the overall Fisher discrimination function remains convex, which greatly improves optimization stability [44].

Compared with LC-KSVD, which introduces label consistency through a learned transformation and classifier, FDDL focuses on improving the statistical distribution of sparse codes. This makes it suitable for classification tasks where the separation between classes is important.

2.7.3 Frozen Dictionary Learning

While most discriminative methods update the entire dictionary during training, certain scenarios benefit from preserving previously acquired knowledge. Frozen dictionary learning, as introduced in the Outlier Learning via Augmented Frozen (OLAF) Dictionaries framework [45], addresses this by fixing a subset of atoms learned from baseline (normal) data while allowing new atoms to adapt to anomalous samples.

The composite dictionary takes the partitioned form $D = [D_f \mid D_u]$, where D_f remains frozen after initial training on normal data, and only the unfrozen portion D_u is updated on data containing deviations. Both K-SVD and Alternating Minimization algorithms can be modified straightforwardly to respect this constraint by skipping updates on frozen columns during the dictionary refinement stage [45].

This approach is particularly useful when preserving existing knowledge is important. In medical imaging, a dictionary learned from normal anatomy can be kept fixed, while additional atoms are learned from abnormal cases. This separation helps maintain normal structural information and improves the interpretability of disease related atoms. However, the effectiveness of this method depends on the quality of the initial frozen dictionary, since an incomplete initial representation may limit the ability of the model to adapt to new patterns.

2.8 Dictionary Learning in Medical Imaging

Dictionary learning has become a powerful tool across the medical imaging pipeline, offering data-adaptive alternatives to fixed transforms such as DCT or wavelets by learning sparse representations directly from image patches [46]. We have identified that DL-based methods are applied to medical imaging in 5 key aspects:

1. Signal Denoising
2. Image Reconstruction
3. Super-Resolution
4. Lesion Detection and Classification
5. Segmentation

2.8.1 Signal Denoising

In denoising, DL-based methods learn a dictionary of atoms from the noisy image itself, then reconstruct a clean image by enforcing that each local patch be sparsely represented by these learnt atoms. This approach has been applied to remove Gaussian noise from natural images [47]. In medical imaging, sparse denoising in diffusion MRI has shown to outperform standard DCT/FFT-based denoising by adapting to the specific structure of the data [48]. Shi et al. proposed an artifact suppressed dictionary learning algorithm (ASDL) [49], which exploited scale and orientation to build discriminative dictionaries for artifact suppression in LDCT.

2.8.2 Image Reconstruction and Super-Resolution

In reconstruction, DL has addressed the inverse problem of recovering images from undersampled k-space acquisitions in compressed sensing MRI [50]. Patch-based dictionaries provide a sparser and more flexible representation than global transforms like wavelets or temporal Fourier bases, enabling more accurate recovery from highly undersampled data.

Closely related to reconstruction is super-resolution, where coupled high-resolution and low-resolution dictionaries are co-trained so that corresponding patches share the same sparse code. This allows high-resolution image patches to be hallucinated from low-resolution inputs. This has been used to enhance cardiac MRI slabs and single-image brain MRI without requiring accurate subpixel alignment [51].

2.8.3 Detection, Classification and Segmentation

Since standard reconstructive dictionaries lack discriminative power, DDL techniques can be used for detection, classification and segmentation. DDL incorporates class labels during training, either by learning separate class-specific dictionaries and using representation residuals for classification, or by jointly learning a shared dictionary alongside a linear classifier over the sparse coding coefficients [46].

Class-specific dictionary learning has been directly applied to CAD for cancer detection. Liu et al. [52] proposed a sparse representation-based classification framework for colorectal polyp and lung nodule detection. The approach learned separate overcomplete dictionaries per class via K-SVD and classifying test samples either by the class whose dictionary yields the most non-zero coefficients when pooled into a global dictionary, or by the class-specific dictionary that achieves the sparsest reconstruction. This generative approach naturally extended to multi-class problems and, when fused with a relevance vector machine baseline through a gated decision tree, achieved consistent sensitivity gains at clinically acceptable false-positive rates on large multi-site colon and lung datasets.

Wu et al. [53] extended class-specific dictionary learning to pulmonary nodule detection in CT, addressing the practical problem that accurate nodule segmentation is often unreliable due to irregular shapes and low contrast. Rather than relying on precisely segmented ROIs, they learned two classification-oriented sub-dictionaries (nodule and non-nodule) together with a shared background dictionary, following the DL-COPAR formulation [54]. Sparse coefficients from the class-specific sub-dictionaries were pooled into feature vectors, refined using DX-score-based feature selection, and

passed to an SVM classifier, yielding good classification performance on the LIDC-IDRI dataset without requiring accurate nodule segmentation, and demonstrating that discriminative dictionary learning can substitute for precise localization as a prerequisite for feature extraction.

Ben-Cohen et al. [55] extended dictionary learning-based approaches from classification into a full lesion detection pipeline, combining a global fully convolutional network (FCN) that generates a lesion probability map across the CT volume with a local, discriminative dictionary-based module for false-positive reduction. Candidate regions from the FCN were partitioned into SLIC superpixels, each represented by a 125-dimensional textural feature vector, and classified using LC-KSVD. A further innovation was online, per-case fine-tuning of the dictionary and linear classifier using high and low-confidence superpixels drawn from each test case’s own FCN probability map, personalizing the discriminative dictionary without additional manual labels. On a clinical liver metastases dataset, this combined FCN-plus-sparsity framework achieved 94.6% sensitivity at 2.9 false positives per case, outperforming the FCN alone, U-Net, a superpixel-only sparse classifier, and a patch-based CNN, with particularly strong performance maintained on small lesions.

Al-Shaikhli et al. extended dictionary learning to a unified multi-class classification and multi-label segmentation framework for brain tumors. It is built around two distinct K-SVD dictionaries learned from a single Flair-MRI modality. A feature dictionary, learned from global topological features combined with gray-level co-occurrence texture features, is used for classification of the volume as normal or abnormal. If abnormal, it is further classified by finding which class-specific dictionary yields the sparsest representation of the test features. Separately, coupled data-label dictionaries are learned jointly from voxel-wise gray-scale patches and their corresponding ground-truth segmentation labels, so that each atom in the image dictionary has an associated atom in the label dictionary; at test time, sparse similarity between test patches and the data dictionary retrieves the matching label-dictionary atoms, which supply foreground/background costs for multi-label graph-cut segmentation of necrosis, edema, and enhancing/non-enhancing tumor. This decoupling of global topological/textural features for classification from local voxel-wise coupled dictionaries for segmentation, using only a single MRI channel, achieved 94.26% classification accuracy and segmentation performance competitive with or exceeding multi-channel BraTS-2013 methods requiring four MRI modalities.

Anitha’s more recent hybrid framework similarly separates dictionary-based segmentation from a distinct classification stage, but replaces reconstruction-residual dictionary classification with a downstream regression-based classifier and shifts from pure MRI to fused 3D MRI-CT input: a Sparse Mahalanobis Dictionary learns class-discriminative atoms for tumor versus non-tumor tissue using a Mahalanobis-distance nonlinearity rather than plain Euclidean sparse coding. These dictionary-derived features feed a Multimodal Markov Random Forest that labels mutually exclusive and exhaustive tissue regions to avoid the tissue loss common in multi-decomposition fusion pipelines. The segmented tumor regions are then classified further by extracting nested 3x3 patches, computing Spearman-correlation-based co-occurrence features, and passing them to a Lagrangian neural network. On combined CT-hemorrhage and MRI tumor datasets this pipeline reported 99.7% accuracy, 98.9% sensitivity, and 98.7% specificity, outperforming CNN, CapsNet, ELM, and SVM-based baselines.

Wang et al. extended coupled dictionary learning to a multi-modality setting for nasopharyngeal carcinoma (NPC) segmentation, addressing the practical constraint that clinical datasets often mix multi-modality patients (with both CT and MRI) and single-modality patients. Building on local coordinate coding, their method jointly learns three coupled dictionaries (CT, MRI, and label) by sharing sparse coefficients across CT and MRI patches for multi-modality samples while separately fitting single-modality patches to their respective dictionary-label pairs, allowing all available data to contribute to a common sparse-coding space rather than discarding incomplete cases. At test time, a voxel is labeled by jointly encoding its CT and MRI patches against the learned CT/MRI dictionaries and projecting the resulting sparse coefficients onto the label dictionary. This joint approach outperformed single-modality dictionary learning (DSC of 76.87% vs. 65.36% for CT-only and 74.94% for MRI-only) and achieved segmentation accuracy comparable to prior specialized NPC methods.

Dong et al. proposed a simultaneous CT reconstruction and segmentation (SRS) framework that uses a discriminative dictionary pair, one for representing image intensity patches and one for representing the corresponding class-probability patches, learned jointly via an extension of K-SVD. The two dictionaries share the same sparse codes for each patch, so that during reconstruction the sparse coding step simultaneously constrains both the intensity reconstruction and the segmentation. The segmentation acts as a built-in regularizer for the otherwise ill-posed inverse problem. Their optimization alternates between sparse coding, an intensity reconstruction update solved via CGLS, and a class-probability update solved with a Frank-Wolfe algorithm, with total-variation regularization on the class probabilities to encourage piecewise-constant, smooth class interfaces. Tested on textured Brodatz-based phantoms and simulated fibrous material samples under varying numbers of projection angles, the method outperformed a non-dictionary SRS baseline and two weaker dictionary-based pipelines, with the gap widening as the number of projections decreased. The authors note the method’s reliance on training data that adequately spans the class-interface geometries expected in the target images, since interfaces not represented in the dictionary tend to be oversmoothed or straightened in the output.

A related work applies DDL to the segmentation problem, once images are already reconstructed. Tong et al. extended their earlier hippocampus-labeling framework into a discriminative dictionary learning for segmentation (DDL_S) method capable of handling multiple structures simultaneously. Their approach jointly learns a reconstructive dictionary and a linear classifier from patches extracted out of a pool of selected atlas images, so that the sparse code of a target patch can be mapped, via the learned classifier, directly to a probabilistic label vector. This produces a subject-specific probabilistic atlas that is then refined with a graph-cuts post-processing step. A central contribution is a voxel-wise local atlas selection strategy (L-DDLS), which selects different subsets of atlases at different spatial locations based on local similarity, rather than relying on one globally selected atlas set (G-DDLS) or a fixed offline-trained dictionary (F-DDLS). This addresses the high inter-subject anatomical variability of abdominal organs, particularly the pancreas, which is poorly served by whole-image similarity measures dominated by large organs like the liver. A multi-resolution Gaussian-pyramid scheme was also incorporated to keep the patch-based approach computationally tractable for large 3D organs. Evaluated on 150 abdominal CT scans, L-DDLS achieved Dice overlaps of 94.9%, 93.6%, 71.1%, and 92.5% for liver, kidneys, pancreas, and spleen respectively, consistently outperforming majority-

voting label fusion, non-local patch-based segmentation, and both G-DDLS and F-DDLS variants, with the largest relative gains for the anatomically variable pancreas. Notably, these results were obtained using only affine (not deformable) registration, underscoring how a learned discriminative dictionary can substitute for some of the burden typically placed on accurate non-rigid registration in multi-atlas segmentation pipelines.

Farhangi’s doctoral thesis develops a shape-based segmentation framework for lung nodules in thoracic CT that relies on a dictionary of previously delineated nodule shapes rather than intensity patches. The method, called Sparse Combination of Training Shapes (SCoTS), embeds nodule boundaries as zero level sets of signed distance functions (SDFs) and builds a dictionary by vectorizing SDFs from a training set of expert-delineated nodules. The evolving segmentation surface is then approximated at each iteration as a sparse linear combination of these dictionary atoms, solved via an l_1 -penalized least-squares formulation using a coordinate-wise “shooting” algorithm. The key contribution is coupling this sparse shape prior with a Chan-Vese region-based active contour. The relative influence of the shape prior versus the low-level image-intensity energy is adjusted dynamically at each iteration based on the sparsity of the reconstruction coefficients, so that as the evolving surface approaches a plausible nodule shape and its representation becomes sparser, the shape prior is trusted more, preventing the contour from leaking into adjacent tissues of similar density. Validated on 542 3D nodules from the LIDC-IDRI database against four independent radiologists, SCoTS achieved a mean DSC of about 0.71-0.72 across radiologist delineations, with the primary failure mode occurring for nodules whose shapes were poorly represented in the training dictionary.

Discriminative dictionary learning has also been extended to multi-modal neuroimaging classification of Alzheimer’s disease (AD) and mild cognitive impairment (MCI). Li et al. built on their previously introduced supervised within-class-similarity discriminative dictionary learning (SCDDL) algorithm by adding a within-class-similarity term to the standard discriminative dictionary objective, which penalizes deviation of each sample’s sparse code from its class mean, alongside the usual reconstruction and linear classification error terms. To handle multiple imaging modalities, they extended SCDDL into a multi-modality version (mSCDDL) that learns a separate dictionary and classifier per modality and fuses per-modality classifier outputs through a weighted combination, with weights optimized by grid search. Using ROI-averaged features from 113 AD, 110 MCI, and 117 normal-control subjects from ADNI, mSCDDL achieved 97.4% accuracy for AD versus normal and 77.7% accuracy for MCI versus normal, outperforming each single-modality SCDDL variant as well as other established multi-modal fusion methods, while requiring substantially less computation time, demonstrating that a compact learned dictionary of just 20 atoms can rival methods that use the entire training set as an implicit dictionary.

The same research group subsequently pushed this framework into a kernelized setting to better exploit cross-modal complementarity. Li et al. introduced multi-feature kernel SCDDL (MKSCDDL), which maps each modality’s features into a higher-dimensional space via Mercer kernels before applying the SCDDL objective, expressing the dictionary implicitly as a linear combination of kernel-mapped training samples. This allows the same reconstruction, classification-error, and within-class-similarity terms can be optimized without explicitly forming the high-dimensional mapping. Sub-kernels for each modality are then combined, extending the earlier

weighted-combination scheme of mSCDDL into a single unified kernel objective. Evaluated on the same three-modality ADNI cohort, MKSCDDL improved classification accuracy over mSCDDL and other multi-modal baselines across all three diagnostic contrasts: 98.2% for AD vs. cognitively unimpaired (CU), 78.5% for MCI vs. CU, and 74.5% for AD vs. MCI, while also reducing test time.

Most relevant to our study, Liyanage et al. tackle COVID-19 detection from chest CT by combining two complementary DL strategies rather than a single dictionary. The first, frozen dictionary learning, addresses class imbalance in the training data, where normal lung-field patches vastly outnumber abnormal patches by learning dictionary atoms hierarchically. The second, LC-KSVD, learns a discriminative dictionary directly from patient-level class labels (normal, common pneumonia, novel coronavirus pneumonia). The two dictionaries are concatenated into a single augmented dictionary, patches are sparse-coded against it via OMP, and simple per-scan statistics (mean, variance, non-zero count) of the sparse coefficients associated with each class-relevant dictionary segment are then fed to an SVM classifier. On the cleaned CC-CCII CT dataset, this patch-based sparse-coding pipeline achieved 89% accuracy for binary NCP-vs-other classification and 82.6% accuracy for the three-class problem, matching or slightly exceeding a benchmark ResNet3D-34 deep learning baseline on the same data. The authors emphasize that because dictionary atoms live in the same domain as the input patches, the resulting classifier offers a more interpretable link between learned features and diagnostic categories than an end-to-end deep network directly echoing the discriminative-dictionary theme seen in Dong et al. and Tong et al., but here applied to whole-scan disease classification rather than pixel-wise segmentation.

Chapter 3

Methodology

Figure 3.1 gives an overview of the full methodology pipeline used in this study. First, the raw CT volumes with their corresponding ground truth segmentation masks are loaded and resampled to a target spacing. Then the CT volumes are preprocessed to isolate the lung region, and their metadata and ground truth masks are used to extract 3D patches of normal and abnormal lung tissue. These patches are then split into training and testing sets, and the training patches are used to learn multiple dictionary learning models with representative atoms for normal lung parenchyma and each abnormal class. Once the dictionaries are learned, they are used to generate sparse codes for each patch in the training set, which are then used to train a linear classifier for multi-class classification. The learned dictionaries are also used to generate the sparse codes for the patches in the test set, which are then classified by the trained linear classifier. The obtained predictions are used to locate the exact voxels in a given patch that contributed to the classification decision, and the generated heatmaps are compared against the ground truth.

3.1 Data Acquisition

The quality and consistency of the input data has a direct bearing on how well the dictionary learning stage can later separate abnormalities from normal lung, so a fair amount of the early project time went into getting this part right before any model code was written. This section walks through the datasets that were chosen, how the




Figure 3.1: Overview of the methodology pipeline.

raw CT volumes were loaded and explored, the different lung segmentation strategies that were tried (and in one case abandoned), and how the final lung volumes were cut down into the 3D patches that feed the K-SVD and LC-KSVD pipelines.

3.1.1 Datasets

Two publicly available chest CT collections were combined for this project rather than relying on a single source, mainly because no single dataset offered both a large enough pool of normal scans and fine-grained, spatially localized annotations for the specific abnormalities of interest.

The first is CT-RATE [23], a large-scale chest CT dataset hosted on Hugging Face that provides volume-level multi-label annotations across 18 abnormality categories. The released label files cover 47,149 training volumes and 3,039 validation volumes, giving 50,188 labelled scans in total. Table-level counts confirm that the dataset is reasonably well populated for the abnormalities most relevant to this study – lung nodules appear as a positive finding in 21,382 scans and consolidation in 8,319 – while also containing 5,226 scans with no positive label at all, which is what made it a practical source of “normal” lung examples. Because of the dataset’s size, only a subset of roughly 2,465 volumes was actually downloaded locally through a small helper script that queries the CT-RATE repository tree via the Hugging Face Hub API and pulls specific volume files on demand.

The second dataset, **ReXGroundingCT** [56], was brought in specifically because it supplies per-finding segmentation masks rather than just volume-level tags. It is organised as 2,992 training, 50 validation and 100 test studies (3,142 in total), each annotated with one or more of 14 finding categories. For this project, attention was narrowed to three categories that best matched the target abnormalities agreed on with the radiologists: atelectasis/consolidation (category 2b), ground-glass opacity (2c) and pulmonary nodules/masses (2d). After filtering the metadata down to studies containing at least one of these categories, 1,418 studies remained (1,354 train, 23 validation, 41 test), split roughly as 659 atelectasis/consolidation, 871 GGO and 908 nodule/mass cases – the two abnormality categories occur together less often than one might expect (correlation of about -0.32), which turned out to be convenient for keeping the classes reasonably separable. A further 347 scans in ReXGroundingCT had no annotated findings at all and were folded in alongside the CT-RATE normals to build a combined pool of negative examples. Each finding mask is stored as a 4D NIfTI volume of shape $[F, H, W, D]$, with one binary channel per annotated finding, which later made it straightforward to collapse multiple findings of the same category into a single lesion mask per scan.

An exploratory data analysis (EDA) was carried out before any model work began, and it shaped a number of decisions described later in this chapter. Beyond simple label counts, the EDA looked at volume dimensions, voxel spacing, and the Hounsfield Unit (HU) intensity distribution inside the lungs for each target category. Raw volumes were consistently 512×512 in-plane, but the number of slices varied widely from scan to scan (values such as 220, 238, 251, 466 and 538 slices were all observed), and voxel spacing was not isotropic either – a typical example measured roughly $0.74\text{ mm} \times 0.74\text{ mm}$ in-plane against 1.5 mm between slices. This kind of inconsistency across scanners and acquisition protocols meant that resampling to a common spacing was unavoidable before patches from different scans could be treated

as comparable inputs, which is discussed further in Section 3.3. Per-category HU histograms computed over lung voxels also showed that nodules, consolidation and GGO occupy overlapping but distinguishable intensity ranges, which gave some early confidence that a patch-based, intensity-driven representation was a reasonable direction to pursue.

3.1.2 Target Abnormality Selection

Rather than picking abnormality classes purely on the basis of which labels happened to have the most examples, the target classes were decided together with practising radiologists over a series of consultation sessions. The datasets on their own contain dozens of possible findings, many of which are either too rare to model reliably or of limited diagnostic priority compared to others, therefore clinical input was needed to narrow this down to something both meaningful and tractable.

The radiologists were shown the available label distributions from CT-RATE and the finding categories in ReXGroundingCT and asked which abnormalities would be most valuable to detect automatically, and which were common enough in practice to justify the modelling effort. Two considerations came up repeatedly in these discussions: first, that pulmonary nodules and ground-glass/consolidation patterns are among the most frequently missed or delayed findings on routine chest CT, and second, that both have a well-established radiological basis for early intervention, making them good candidates for a proof-of-concept detection system. Based on this feedback, and cross-checked against how well each candidate class was actually represented in the annotated data, the following were fixed as the target classes for this study:

- 2C : Ground-glass opacity (GGO) / consolidation
- 2d : Lung nodule

This clinical grounding mattered in a second, more practical way as well: it directly influenced the segmentation and patch-extraction requirements described in the rest of this section, since both target abnormalities tend to sit close to the lung boundary or alter the surrounding tissue in ways that are easy for naive segmentation methods to mishandle.

3.1.3 CT Volume Loading

All CT scans in this study were distributed in NIfTI format, which keeps both the 3D voxel array and the spatial metadata (orientation, spacing, origin) bundled together in one file. Volumes were read using `nibabel` and converted into float32 HU arrays, with each scan represented as

$$I(x, y, z) \in \mathbb{R}^{H \times W \times D} \quad (3.1)$$

where H , W and D denote the height, width and depth of the volume. Alongside the intensity array, the voxel spacing recorded in the NIfTI header was retained for every scan, since, as noted above, spacing was not consistent across the dataset and this information is needed later to resample volumes onto a common grid.

Where ground-truth finding masks were available (ReXGroundingCT), the corresponding 4D mask volume was loaded through the same pathway and kept aligned with its source CT volume via a shared volume identifier, therefore that a mask could always be traced back to the exact scan and finding it belongs to.

Loading a scan or a mask, however, was never just a matter of reading the NIfTI file and handing back an array – both paths funnel through a shared loading routine that also decides whether resampling and lung preprocessing are required before the data is considered ready to use. Figure 3.2 summarises this routine. After a volume or mask is read with NiBabel, its voxel spacing is compared against the target spacing of 1.5 mm isotropic; if the two already match, the array is passed through unchanged, but if not, it is resampled using trilinear interpolation for CT intensities or nearest-neighbour interpolation for masks, mirroring the distinction made earlier for `resample_volume` and `resample_mask` – nearest-neighbour is used specifically to avoid inventing intermediate label values between the discrete finding categories in a mask. Only volumes, and never masks, are then routed through the preprocessing block, where lung segmentation, HU clipping to $[-1000, 500]$ and min-max normalisation are applied in sequence before the final volume is returned; masks skip this block entirely and are returned as soon as resampling is done, since a binary or categorical finding mask has no HU values to clip or normalise and does not need to be run through the lung segmentation model. Keeping this branching logic in one place meant that every scan and every mask used later in the pipeline, regardless of which dataset it came from, had been through exactly the same spacing correction and (where relevant) the same preprocessing steps, so patches extracted from different scans remained directly comparable.




Figure 3.2: Shared loading pipeline for CT volumes and finding masks, showing the spacing check, resampling branch (trilinear for volumes, nearest-neighbour for masks) and the preprocessing block applied only to volumes.

3.2 Preprocessing

Before any patches could be extracted, the raw volumes had to be reduced to just the lung region. A whole-volume approach was ruled out fairly early: bones, soft tissue

and background air all add computational overhead without contributing anything useful to dictionary learning, and worse, they risk letting the model pick up spurious patterns unrelated to lung pathology. Three different approaches to isolating the lungs were tried over the course of the project, in the order described below, before settling on a final method.

3.2.1 Classical Image Processing-Based Lung Segmentation

The first attempt used a fairly standard HU-threshold pipeline, since lung tissue is normally well separated from surrounding structures in intensity. Each volume was first thresholded to keep only voxels between -900 and -200 HU, the range typically associated with aerated lung parenchyma:

$$M_{thresh}(x, y, z) = \begin{cases} 1, & -900 \leq I(x, y, z) \leq -200 \\ 0, & \text{otherwise} \end{cases} \quad (3.2)$$

This crude threshold mask is noisy on its own, so it was cleaned up with a short sequence of morphological steps applied slice-by-slice and volume-wise using `scipy.ndimage` and `skimage.segmentation`. Border-connected regions were removed on each axial slice with `clear_border` (mainly to drop the surrounding air and imaging-table artefacts that also fall inside the HU window), after which only the two largest connected components were kept – under the assumption that these correspond to the left and right lung:

A 3D binary closing (a $3 \times 3 \times 3$ connectivity structure, 2 iterations) was then used to smooth the boundary and reconnect thin gaps, followed by `binary_fill_holes` to plug any internal cavities left inside the mask – these two steps are exactly what a lung mask needs, since vessels, bronchi and small nodules would otherwise show up as holes in an otherwise solid lung region.

On visibly normal scans, this five-stage pipeline (threshold $\rightarrow$ clear border $\rightarrow$ largest components $\rightarrow$ closing $\rightarrow$ hole filling) produced reasonable-looking lung masks when checked slice-by-slice. The problem showed up once pathology entered the picture. Consolidation and areas of lung opacity raise the local tissue density towards that of soft tissue, so the affected voxels are not registered as lung parenchyma at all under the initial -900 to -200 HU window – they get cut out of the mask at the very first thresholding step, before any of the later stages ever see them.

What happens next depends on where the abnormality sits and how large it is. When the affected region is small and fully surrounded by normal, correctly thresholded lung tissue, it behaves like just another gap in the mask, and the binary closing and hole-filling steps are usually enough to patch it back in – the same mechanism that is meant to fill in small vessels and airways ends up rescuing these interior lesions almost by accident. But when the abnormality is larger, or sits close to the lung boundary rather than fully enclosed within it, closing and hole-filling can no longer bridge the gap: the affected region gets left connected to the outside of the lung rather than treated as an enclosed hole, so it is excluded from the final mask along with everything else outside the lung boundary. This is exactly the pattern that shows up when the intermediate masks are stepped through and overlaid on the source slice – a large opacity running along the pleural surface is visibly missing from the final lung segmentation and from the CT after the mask is applied, while the ground-truth annotation for that same region sits entirely outside the segmented

lung. The five-stage pipeline was checked this way across several scans, and the failure was consistent whenever the abnormal region was either large or adjacent to the lung border, rather than an isolated case.

For a project whose entire purpose is to detect exactly these kinds of abnormalities, a segmentation step that only reliably preserves lesions small enough to be treated as incidental holes is not something that could be fixed with more morphological post-processing – the information is already lost at the thresholding step, since abnormal tissue often does not fall inside the normal lung parenchyma HU window in the first place. This is what motivated the move away from purely intensity-based segmentation and towards a learned segmentation model, LungMask’s **LMInferer**, described next.

3.2.2 LungMask-Based Lung Segmentation

The method that was ultimately adopted is the **LungMask** library, accessed through its **LMInferer** interface. Unlike the classical pipeline, LungMask is a trained segmentation network built specifically for whole-lung extraction, and it proved noticeably more stable when abnormal tissue was present – consolidation and GGO regions inside the lung boundary were retained rather than excluded, which directly addressed the main failure observed with classical image processing approach.

Given a CT volume, LungMask returns a label map with three values: 0 for background, 1 for the left lung and 2 for the right lung. This was converted into a binary mask and applied to the original volume:

$$I_{lung} = I \times M_{lung} \quad (3.3)$$

where I is the original CT volume and M_{lung} is the binary lung mask produced by LungMask. Voxels falling outside the mask were reset to a fixed background value of -1000 HU (approximating air), after which the volume was clipped to the range $[-1000, 500]$ HU and linearly rescaled to $[0, 1]$. This combination of segmentation, HU clipping and normalisation was wrapped into a single preprocessing routine that also resamples each volume to an isotropic 1.5 mm spacing beforehand, therefore that every scan reaching the patch extraction stage has a consistent voxel grid regardless of its original acquisition parameters.

3.2.3 Patch Extraction

With a clean, normalised lung volume in hand, the final data acquisition step was to break each volume down into smaller 3D patches. Feeding entire volumes directly into the dictionary learning stage was never really practical – a single scan can contain tens of millions of voxels – and patch-based sampling has the added benefit of letting the model focus on local structural patterns rather than whole-organ context, which fits naturally with how K-SVD and LC-KSVD represent signals as sparse combinations of small dictionary atoms.

Each patch is a fixed $16 \times 16 \times 16$ voxel cube, which at the resampled 1.5 mm spacing corresponds to roughly a 24 mm^3 region of lung tissue – large enough to capture a typical small nodule or a local patch of opacity, without dragging in unrelated anatomy. Patches were sampled differently depending on whether they were meant to represent normal tissue or a target abnormality:

1. **Normal patches** were drawn on a non-overlapping grid across the entire lung volume, with a stride equal to the patch size. A patch was discarded if more than half of its voxels were near-zero after masking, since that generally indicates the patch sits mostly outside the lung region rather than within actual tissue.
2. **Abnormal patches** were sampled only from scans carrying a ground-truth finding mask. Since a single scan can have several individual findings that belong to the same target category, the relevant finding channels of the 4D mask were first OR-combined into one binary lesion mask per category present in that scan. A tight axis-aligned bounding box was then computed around the foreground of that lesion mask, and the patch window was slid across the box with a stride of 4 voxels in each direction – markedly denser than the normal-patch grid, since the region of interest here is already small and every overlapping view of it is potentially useful. Each candidate patch was only kept if the fraction of lesion-labelled voxels within its corresponding mask region reached a category-specific threshold: 50% for ground-glass opacity/consolidation patches, but only 10% for nodule/mass patches, reflecting how much smaller and more localised a pulmonary nodule typically is compared with a spread-out opacity – a fixed 50% cutoff would have discarded almost every genuine nodule patch. A scan positive for both target categories simply contributes patches to both, each filtered against its own bounding box and threshold.

Figure 3.3 lays out this branching more concretely. The first decision is simply whether the loaded scan is one of the designated normal volumes or one that carries a ground-truth abnormality mask. Normal scans go straight into a regular sampling grid built over the whole volume, whereas scans with an annotation instead have a 3D bounding box generated directly around the ground-truth lesion, so that patch sampling is concentrated on the region that actually contains the finding rather than scattered across the entire lung. Whichever route a patch takes, it still has to pass one final check before it is accepted: patches coming from the normal grid are rejected if more than half of their voxels sit at background level (-1000 HU, i.e. outside the lung after masking), while patches coming from a lesion bounding box are rejected unless their lesion-voxel fraction meets that category’s threshold. Only patches that clear this check are appended to the patch matrix; everything else is discarded before it ever reaches the dictionary learning stage. Structuring the two branches this way keeps the normal and abnormal sampling logic independent while funnelling both into the same final acceptance criterion and the same output matrix.

Formally, each extracted patch can be written as

$$P_i \in \mathbb{R}^{p \times p \times p}, \quad p = 16 \quad (3.4)$$

and is subsequently flattened into a vector representation before being handed to the dictionary learning algorithms:

$$x_i \in \mathbb{R}^n, \quad n = p^3 = 4096 \quad (3.5)$$

Patches extracted this way were saved per scan and per class (normal, GGO, nodule) as compressed archives containing the flattened feature matrix, class labels, the originating scan identifiers and the voxel coordinates of each patch, and were later concatenated across classes to build the full training matrix used for K-SVD,




Figure 3.3: Patch extraction workflow, showing the normal-scan sampling grid and the lesion bounding-box branch for annotated scans, each filtered by its own background/lesion-fraction threshold before patches are added to the patch matrix.

frozen dictionary learning, and LC-KSVD. This patch-based representation keeps memory requirements manageable while still allowing the learned dictionary atoms to specialise in the local structural signatures associated with each target pulmonary abnormality.

3.3 Evaluation Metrics

The evaluation of medical image analysis systems requires multiple metrics because different measures capture different aspects of model behaviour. In chest CT analysis, classification metrics describe the ability of the model to identify abnormal findings, segmentation metrics measure spatial agreement with reference annotations, and interpretability metrics evaluate whether the learned representations provide meaningful explanations.

3.3.1 Classification

For a binary classification problem, let TP , TN , FP , and FN represent true positives, true negatives, false positives, and false negatives, respectively. Accuracy measures the overall proportion of correctly classified samples:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}. \quad (3.6)$$

Although accuracy provides an overall performance estimate, it can become misleading when the dataset contains class imbalance. Therefore, additional metrics based on positive class prediction are considered.

Precision indicates the reliability of positive predictions by measuring the fraction of detected abnormalities that are correct:

$$\text{Precision} = \frac{TP}{TP + FP}. \quad (3.7)$$

Recall, also referred to as sensitivity, measures the ability of the model to recover existing abnormal cases:

$$\text{Recall} = \frac{TP}{TP + FN}. \quad (3.8)$$

The F1-score combines precision and recall into a single measure and provides a balanced evaluation when both false positive and false negative errors are important:

$$\text{F1} = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}. \quad (3.9)$$

The Area Under the Receiver Operating Characteristic Curve (AUROC) evaluates the ability of a classifier to separate positive and negative samples over all possible decision thresholds. It is calculated from the relationship between the true positive rate (TPR) and false positive rate (FPR):

$$\text{AUROC} = \int_0^1 \text{TPR}(\text{FPR}) d\text{FPR}. \quad (3.10)$$

A higher AUROC indicates that the model can better distinguish between abnormal and normal cases without depending on a particular classification threshold.

Average Precision (AP) summarizes the precision-recall relationship and is particularly useful for medical datasets where abnormal cases may represent a smaller proportion of the data:

$$\text{AP} = \sum_n (R_n - R_{n-1}) P_n, \quad (3.11)$$

where P_n and R_n represent precision and recall values at different classification thresholds. Higher average precision indicates better detection of positive cases under class imbalance.

The reliability of predicted probabilities is measured using Expected Calibration Error (ECE). The metric compares predicted confidence with the actual accuracy within confidence intervals:

$$\text{ECE} = \sum_{m=1}^M \frac{|B_m|}{N} |\text{acc}(B_m) - \text{conf}(B_m)|. \quad (3.12)$$

Lower ECE values indicate that the predicted confidence scores better reflect the actual likelihood of correctness.

3.3.2 Localization and Segmentation

For segmentation evaluation, let P represent the predicted region and G represent the ground-truth annotation. The Dice coefficient measures the overlap between the two regions:

$$\text{Dice}(P, G) = \frac{2|P \cap G|}{|P| + |G|}. \quad (3.13)$$

Intersection over Union (IoU), also known as the Jaccard index, measures the ratio between the intersection and union of the predicted and reference regions:

$$\text{IoU}(P, G) = \frac{|P \cap G|}{|P \cup G|}. \quad (3.14)$$

Dice and IoU provide complementary information about segmentation quality, where higher values indicate better spatial agreement between predicted and reference regions.

3.3.3 Interpretability of Sparse Dictionary Learning

In sparse dictionary learning, interpretability is obtained from the learned dictionary atoms and their corresponding sparse representations. Therefore, evaluation focuses on whether the dictionary provides compact, meaningful, and discriminative representations of CT image patterns.

The reconstruction error measures how effectively the learned dictionary represents the input data. Given an input patch x , dictionary D , and sparse coefficient vector α , the reconstruction error is defined as:

$$\text{Reconstruction Error} = \|x - D\alpha\|_2^2. \quad (3.15)$$

A lower reconstruction error indicates that the dictionary captures important structures present in the image data.

The sparsity of the representation describes how many atoms contribute to the reconstruction. It is measured using the number of non-zero coefficients:

$$\text{Sparsity} = \|\alpha\|_0. \quad (3.16)$$

Lower values indicate that only a small number of atoms are required to explain a sample, making the representation easier to interpret.

Dictionary coherence evaluates the similarity between different atoms in the learned dictionary. The mutual coherence is defined as:

$$\mu(D) = \max_{i \neq j} \frac{|d_i^T d_j|}{\|d_i\|_2 \|d_j\|_2}. \quad (3.17)$$

A lower coherence value indicates that the dictionary contains less redundant atoms and that individual atoms capture more distinct image characteristics.

Atom utilization measures how frequently each dictionary atom contributes to the representation of the dataset:

$$U_k = \sum_{i=1}^N \mathbb{I}(\alpha_{k,i} \neq 0), \quad (3.18)$$

where U_k represents the activation frequency of atom k . This metric helps identify whether the learned dictionary is effectively utilized or dominated by a limited number of atoms.

For supervised dictionary learning approaches such as LC-KSVD, class-specific atom activation provides additional interpretability. Disease-related atoms are expected to activate more frequently for samples belonging to their associated class compared with unrelated classes. This indicates that the learned atoms capture discriminative patterns rather than only general image structures.

Finally, visual inspection of reconstructed dictionary atoms is used as a qualitative interpretability assessment. By examining representative atoms, their relationship with anatomical structures or pathological appearances can be analysed, providing an intuitive understanding of the features learned by the model.

Chapter 4

Conclusion

Appendix A

Essential English Grammar - A Handout

The Appendices are numbered as A, B,.. and so on. Sections and Subsections of the appendix then would be A.1, A.1.1, A.2 etc.

A.1 Simple Present Tense

A.2 Perfect Tense

A.3 Past Past Perfect Tense

A.4 Passive Voice

A.5 Technical Writing Fundamentals

A.6 Proof Reading

Appendix B

Graphical Interpretation of Data

The graphical interpretation given below shows the relationship between x and y is acceptable to the purpose of the project

B.1 Data Analysis

B.1.1 Statistical Analysis

B.1.2 Empirical Analysis

B.2 The X vs Y vs Z

Bibliography

- [1] C. J. Murray, “The global burden of disease study at 30 years,” *Nature medicine*, vol. 28, no. 10, pp. 2019–2026, 2022.
- [2] K. Doi, “Computer-aided diagnosis in medical imaging: historical review, current status and future potential,” *Computerized medical imaging and graphics*, vol. 31, no. 4-5, pp. 198–211, 2007.
- [3] A. Borakati, A. Perera, J. Johnson, and T. Sood, “Diagnostic accuracy of x-ray versus ct in covid-19: a propensity-matched database study,” *BMJ open*, vol. 10, no. 11, p. e042946, 2020.
- [4] G. McCreadie and T. Oliver, “Eight ct lessons that we learned the hard way: an analysis of current patterns of radiological error and discrepancy with particular emphasis on ct,” *Clinical radiology*, vol. 64, no. 5, pp. 491–499, 2009.
- [5] M. Ennab and H. Mcheick, “Advancing ai interpretability in medical imaging: a comparative analysis of pixel-level interpretability and grad-cam models,” *Machine Learning and Knowledge Extraction*, vol. 7, no. 1, p. 12, 2025.
- [6] N. L. S. T. R. Team, “The national lung screening trial: overview and study design,” *Radiology*, vol. 258, no. 1, pp. 243–253, 2011.
- [7] H. J. de Koning, C. M. van Der Aalst, P. A. de Jong, E. T. Scholten, K. Nackaerts, M. A. Heuvelmans, J.-W. J. Lammers, C. Weenink, U. Yousaf-Khan, N. Horeweg, *et al.*, “Reduced lung-cancer mortality with volume ct screening in a randomized trial,” *New England journal of medicine*, vol. 382, no. 6, pp. 503–513, 2020.
- [8] P. J. Withers, C. Bouman, S. Carmignato, V. Cnudde, D. Grimaldi, C. K. Hagen, E. Maire, M. Manley, A. Du Plessis, and S. R. Stock, “X-ray computed tomography,” *Nature Reviews Methods Primers*, vol. 1, no. 1, p. 18, 2021.
- [9] K. Doi, “Computer-aided diagnosis in medical imaging: Historical review, current status and future potential,” *Computerized Medical Imaging and Graphics*, vol. 31, no. 4, pp. 198–211, 2007. Computer-aided Diagnosis (CAD) and Image-guided Decision Support.
- [10] T. E. Cupples, J. E. Cunningham, and J. C. Reynolds, “Impact of computer-aided detection in a regional screening mammography program,” *American Journal of Roentgenology*, vol. 185, no. 4, pp. 944–950, 2005. PMID: 16177413.
- [11] T. W. Freer and M. J. Ulissey, “Screening mammography with computer-aided detection: Prospective study of 12,860 patients in a community breast center,” *Radiology*, vol. 220, no. 3, pp. 781–786, 2001. PMID: 11526282.

- [12] F. Li, M. Aoyama, J. Shiraishi, H. Abe, Q. Li, K. Suzuki, R. Engelmann, S. Sone, H. Macmahon, and K. Doi, “Radiologists’ performance for differentiating benign from malignant lung nodules on high-resolution ct using computer-estimated likelihood of malignancy,” *AJR. American journal of roentgenology*, vol. 183, pp. 1209–15, 12 2004.
- [13] T. Zhao, Y. Gu, J. Yang, N. Usuyama, H. H. Lee, S. Kiblawi, T. Naumann, J. Gao, A. Crabtree, J. Abel, *et al.*, “A foundation model for joint segmentation, detection and recognition of biomedical objects across nine modalities,” *Nature methods*, vol. 22, no. 1, pp. 166–176, 2025.
- [14] J. Wasserthal, H. C. Breit, M. T. Meyer, *et al.*, “Totalsegmentator: Robust segmentation of 104 anatomic structures in ct images,” *Radiology: Artificial Intelligence*, vol. 5, no. 5, p. e230024, 2023.
- [15] Z. Yang, Z. Chen, Y. Sun, A. Strittmatter, A. Raj, A. Allababidi, J. S. Rink, and F. G. Zöllner, “seg2med: a segmentation-based medical image generation framework using denoising diffusion probabilistic models,” *arXiv e-prints*, pp. arXiv–2504, 2025.
- [16] L. Blankemeier, A. Kumar, J. P. Cohen, J. Liu, L. Liu, D. Van Veen, S. J. S. Gardezi, H. Yu, M. Paschali, Z. Chen, J.-B. Delbrouck, E. Reis, R. Holland, C. Truyts, C. Bluethgen, Y. Wu, L. Lian, M. E. K. Jensen, S. Ostmeier, M. Varma, J. M. J. Valanarasu, Z. Fang, Z. Huo, Z. Nabulsi, D. Ardila, W.-H. Weng, E. A. Junior, N. Ahuja, J. Fries, N. H. Shah, G. Zaharchuk, M. Willis, A. Yala, A. Johnston, R. D. Boutin, A. Wentland, C. P. Langlotz, J. Hom, S. Gatidis, and A. S. Chaudhari, “Merlin: a computed tomography vision–language foundation model and dataset,” *Nature*, vol. 652, p. 1318–1328, Mar. 2026.
- [17] J. Ma, Y. He, F. Li, L. Han, C. You, and B. Wang, “Segment anything in medical images,” *Nature Communications*, vol. 15, p. 654, 2024.
- [18] N. Ravi, J. Ma, Y. He, B. Wang, *et al.*, “Medsam2: Segment anything in 3d medical images and videos,” *arXiv preprint arXiv:2504.03600*, 2025.
- [19] E. M. Ayllon *et al.*, “Can foundation models really segment tumors? a benchmarking odyssey in lung ct imaging,” *arXiv preprint arXiv:2505.01239*, 2025.
- [20] Meta AI, “Sam 3: Segment anything with concepts.” Submitted to ICLR 2026 (under review), 2025.
- [21] A. Liu, R. Xue, X. R. Cao, Y. Shen, Y. Lu, X. Li, Q. Chen, and J. Chen, “Medsam3: Delving into segment anything with medical concepts,” *arXiv preprint arXiv:2511.19046*, 2025.
- [22] I. E. Hamamci, S. Er, *et al.*, “Ct-clip: Contrastive vision-language pretraining for chest ct,” *arXiv preprint arXiv:2403.17834*, 2024.
- [23] I. E. Hamamci, S. Er, E. Almas, E. Keles, E. Keles, M. Basar, A. Gungor, and G. Unal, “Ct-rate: A large-scale dataset for chest ct radiology report generation and abnormality detection,” *arXiv preprint arXiv:2401.07515*, 2024.

- [24] E. Tjoa and C. Guan, “A survey on explainable artificial intelligence (xai): Toward medical xai,” *IEEE Transactions on Neural Networks and Learning Systems*, vol. 32, no. 11, pp. 4793–4813, 2021.
- [25] M. T. Ribeiro, S. Singh, and C. Guestrin, “Why should i trust you?: Explaining the predictions of any classifier,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 1135–1144, ACM, 2016.
- [26] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” *Advances in neural information processing systems*, vol. 30, 2017.
- [27] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, “Grad-cam: Visual explanations from deep networks via gradient-based localization,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pp. 618–626, 2017.
- [28] A. Chattopadhyay, A. Sarkar, P. Howlader, and V. N. Balasubramanian, “Grad-cam++: Generalized gradient-based visual explanations for deep convolutional networks,” in *2018 IEEE Winter Conference on Applications of Computer Vision (WACV)*, pp. 839–847, IEEE, 2018.
- [29] H. Wang, Z. Wang, M. Du, F. Yang, Z. Zhang, S. Ding, P. Mardziel, and X. Hu, “Score-cam: Score-weighted visual explanations for convolutional neural networks,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pp. 111–119, 2020.
- [30] P.-T. Jiang, C.-B. Zhang, Q. Hou, M.-M. Cheng, and Y. Wei, “Layercam: Exploring hierarchical class activation maps for localization,” *IEEE transactions on image processing*, vol. 30, pp. 5875–5888, 2021.
- [31] V. Hassija, V. Chamola, A. Mahapatra, A. Singal, D. Goel, K. Huang, S. Scardapane, I. Spinelli, M. Mahmud, and A. Hussain, “Interpreting black-box models: a review on explainable artificial intelligence,” *Cognitive Computation*, vol. 16, no. 1, pp. 45–74, 2024.
- [32] L. Ribeiro-Flucht, X. Chen, and D. Meurers, “Explainable ai in language learning: Linking empirical evidence and theoretical concepts in proficiency and readability modeling of portuguese,” in *Proceedings of the 19th workshop on innovative use of NLP for building educational applications (BEA 2024)*, pp. 199–209, 2024.
- [33] F. Doshi-Velez and B. Kim, “Towards a rigorous science of interpretable machine learning,” *arXiv preprint arXiv:1702.08608*, 2017.
- [34] R. Guidotti, “Counterfactual explanations and how to find them: literature review and benchmarking: R. guidotti,” *Data Mining and Knowledge Discovery*, vol. 38, no. 5, pp. 2770–2824, 2024.
- [35] M. Van Lent, W. Fisher, and M. Mancuso, “An explainable artificial intelligence system for small-unit tactical behavior,” in *Proceedings of the national conference on artificial intelligence*, pp. 900–907, Menlo Park, CA; Cambridge, MA; London; AAAI Press; MIT Press; 1999, 2004.

- [36] H. Zheng, H. Yong, and L. Zhang, “Deep convolutional dictionary learning for image denoising,” in *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 630–641, 2021.
- [37] Y. Li, S. Ding, and Z. Li, “Dictionary learning with the cospase analysis model based on summation of blocked determinants as the sparseness measure,” *Digital Signal Processing*, vol. 48, pp. 298–309, 2016.
- [38] H. Lee, A. Battle, R. Raina, and A. Ng, “Efficient sparse coding algorithms,” in *Advances in Neural Information Processing Systems* (B. Schölkopf, J. Platt, and T. Hoffman, eds.), vol. 19, MIT Press, 2006.
- [39] M. Aharon, M. Elad, and A. Bruckstein, “K-svd: An algorithm for designing overcomplete dictionaries for sparse representation,” *IEEE Transactions on Signal Processing*, vol. 54, no. 11, pp. 4311–4322, 2006.
- [40] A. Yahyaoui, N. Fnaiech, and F. Fnaiech, “A suitable initialization procedure for speeding a neural network job-shop scheduling,” *IEEE Transactions on Industrial Electronics*, vol. 58, no. 3, pp. 1052–1060, 2011.
- [41] Y. Pati, R. Rezaiifar, and P. Krishnaprasad, “Orthogonal matching pursuit: recursive function approximation with applications to wavelet decomposition,” in *Proceedings of 27th Asilomar Conference on Signals, Systems and Computers*, pp. 40–44 vol.1, 1993.
- [42] M. Aharon, M. Elad, and A. Bruckstein, “K-svd: An algorithm for designing overcomplete dictionaries for sparse representation,” *IEEE Transactions on Signal Processing*, vol. 54, no. 11, pp. 4311–4322, 2006.
- [43] Z. Jiang, Z. Lin, and L. S. Davis, “Label consistent k-svd: Learning a discriminative dictionary for recognition,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 35, no. 11, pp. 2651–2664, 2013.
- [44] M. Yang, L. Zhang, X. Feng, and D. Zhang, “Fisher discrimination dictionary learning for sparse representation,” *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pp. 543–550, 2011.
- [45] B. T. Carroll, B. M. Whitaker, W. Dayley, and D. V. Anderson, “Outlier learning via augmented frozen dictionaries,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 25, no. 6, pp. 1207–1219, 2017.
- [46] T. Tong, J. Caballero, K. Bhatia, and D. Rueckert, “Dictionary learning for medical image denoising, reconstruction, and segmentation,” in *Machine learning and medical imaging*, pp. 153–181, Elsevier, 2016.
- [47] M. Elad and M. Aharon, “Image denoising via sparse and redundant representations over learned dictionaries,” *IEEE Transactions on Image processing*, vol. 15, no. 12, pp. 3736–3745, 2006.
- [48] L. Bao, M. Robini, W. Liu, and Y. Zhu, “Structure-adaptive sparse denoising for diffusion-tensor mri,” *Medical image analysis*, vol. 17, no. 4, pp. 442–457, 2013.

- [49] L. Shi, Y. Chen, H. Shu, L. Luo, C. Toumoulin, and J.-L. Coatrieux, “Low-dose ct image processing using artifact suppressed dictionary learning,” in *2014 IEEE 11th International Symposium on Biomedical Imaging (ISBI)*, pp. 1127–1130, IEEE, 2014.
- [50] S. Ravishankar and Y. Bresler, “Mr image reconstruction from highly under-sampled k-space data by dictionary learning,” *IEEE Transactions on Medical Imaging*, vol. 30, no. 5, pp. 1028–1041, 2011.
- [51] A. Rueda, N. Malpica, and E. Romero, “Single-image super-resolution of brain mr images using overcomplete dictionaries,” *Medical image analysis*, vol. 17, no. 1, pp. 113–132, 2013.
- [52] M. Liu, L. Lu, X. Ye, S. Yu, and M. Salganicoff, “Sparse classification for computer aided diagnosis using learned dictionaries,” in *International Conference on Medical Image Computing and Computer-Assisted Intervention*, pp. 41–48, Springer, 2011.
- [53] P. Wu, K. Xia, Y. Zhang, X. Qian, G. Wang, and H. Yu, “Dictionary learning-based ct detection of pulmonary nodules,” in *Developments in X-Ray Tomography X*, vol. 9967, pp. 88–99, SPIE, 2016.
- [54] S. Kong and D. Wang, “A dictionary learning approach for classification: Separating the particularity and the commonality,” in *European conference on computer vision*, pp. 186–199, Springer, 2012.
- [55] A. Ben-Cohen, E. Klang, A. Kerpel, E. Konen, M. M. Amitai, and H. Greenspan, “Fully convolutional network and sparsity-based dictionary learning for liver lesion detection in ct examinations,” *Neurocomputing*, vol. 275, pp. 1585–1594, 2018.
- [56] P. Rajpurkar *et al.*, “Rexgroundingct: A large-scale dataset for grounded radiology report generation on chest ct,” *arXiv preprint*, 2024.