



An interpretable Dictionary Learning approach for Diagnostic Hypotheses from 3D Chest

An undergraduate project report submitted to the

Department of Electrical and Information Engineering
Faculty of Engineering
University of Ruhuna
Sri Lanka

in partial fulfillment of the requirements for the

**Degree of the Bachelor of the Science of Engineering
Honours**

by

K.R.C.S. Kekulawala - EG/2021/4609
M.D.Y.B. Kularathne - EG/2021/4620
B.G.P.U. Thilakshana - EG/2021/4831
B.P.G.L. Thejana - EG/2021/4828

11th May 2026

.....
Dr. Kaveen Liyanage
(Supervisor)

.....
Mrs. Sithara Mahagama
(Co-Supervisor)

Abstract

Multi-class abnormality detection from 3D chest Computed Tomography requires clinically interpretable methods to support radiological decision-making. Current state-of-the-art (SotA) models often rely on post-hoc explainability techniques that provide approximate localizations and lack direct traceability to the models decisions. Furthermore, SotA models rarely support simultaneous multi-class classification and grounded localization, leading to a fragmented approach to diagnosis. This study aims to address these challenges through two contributions: First, we introduce an inherently interpretable sparse dictionary learning framework in which detected abnormalities, such as Pulmonary Nodules, Ground-Glass Opacity and Consolidation, can be back-projected to the original voxel space, allowing diagnostic decisions to be traced back to relevant anatomical structures. Second, we develop a unified evaluation framework to benchmark SotA models across three axes: (i) classification accuracy, (ii) segmentation quality, and (iii) explainability, using standardized metrics and a publically available 3D chest CT dataset with expertly annotated abnormality localizations. Our method is benchmarked against SotA models using this framework, demonstrating competitive classification and segmentation performance while providing improved interpretability compared to deep learning baselines.

Acknowledgments

We would like to express our sincere gratitude to Dr. Kaveen Liyanage (Supervisor), Mrs. Sithara Mahagama (Co-Supervisor), and Dr. Subhodini Indrapala (Consultant Radiologist at Teaching Hospital, Karapitiya), for their invaluable guidance, support, and constructive feedback throughout the course of this project. Their expertise, mentorship, and encouragement have been instrumental in shaping the direction and quality of this research.

We are grateful to the Department of Electrical and Information Engineering for providing access to computing resources and research facilities that were essential for conducting this project. The department’s commitment to fostering an environment conducive to research and innovation has been greatly appreciated.

We would like to acknowledge the creators and maintainers of the publicly available datasets used in this research, particularly the CT-RATE and RexGrounding-CT datasets and other medical imaging resources that have made this work possible. These datasets have been crucial for developing and validating the proposed methods.

We extend our thanks to the developers of the open-source frameworks and libraries used in this project, including PyTorch, scikit-image, and other tools that significantly contributed to the implementation of this research.

Contents

Abstract	ii
Acknowledgements	iii
Contents	v
List of Figures	vi
List of Tables	vii
Acronyms	viii
1 Introduction	ix
1.1 Background	ix
1.2 Problem Statement	ix
1.3 Objectives and Scope	x
1.3.1 Objectives	x
1.3.2 Scope	x
1.4 Literature Review	x
1.4.1 Classification Architectures	xi
1.4.2 Segmentation Architectures	xi
1.4.3 Unified Architectures	xii
1.4.4 Dictionary Learning Approaches	xii
1.4.5 Explainability Approaches	xii
1.5 Report Outline	1
2 Background	2
2.1 Dictionary Learning	2
2.2 Radiological Patterns in Lung CT Imaging	3
2.2.1 CT Imaging	3
2.2.2 Ground Glass Opacity (GGO) and Consolidation	3
2.2.3 Pulmonary Nodules	3
2.3 Vision Transformers (ViTs)	4
3 Methodology	5
3.1 Data Acquisition and Preprocessing	5
3.2 Implementation of the LC-KSVD Approach	6
3.3 Abnormality Classification	6
3.4 Back-Projecting Class-Discriminative Atoms	6
3.5 Deriving Explainability through Segmentation Masks	7

4	Evaluation Framework and Results	8
4.1	Evaluation Framework Design	8
4.2	Evaluation Metrics	8
4.2.1	Classification Metrics	9
4.2.2	Localization and Grounding Metrics	9
4.2.3	Interpretability Metrics	10
4.3	Classification Accuracy Evaluation	10
4.3.1	Overall Classification Performance	11
4.3.2	Per-Class Analysis	11
4.4	Grounded Localization Evaluation	12
4.4.1	Per-Class Localization Performance	12
4.4.2	Qualitative Comparison of Localization Outcomes	13
4.5	Inherent Explainability Evaluation	14
4.5.1	Atom-to-Voxel Traceability	14
4.5.2	Spatial Faithfulness of Explanations	14
4.6	Summary of Evaluation Findings	15
5	Discussion	16
A	Appendix 1	17
B	Appendix 2	18
	APPENDIX	18
	Bibliography	19

List of Figures

4.1	Architecture diagram of the proposed evaluation framework.	9
4.2	Overall classification performance comparison across models.	11
4.3	Per-class AUROC comparison Across Models.	11
4.4	Overall localization performance comparison for BiomedParse and Merlin.	12
4.5	Mean Dice comparison by lesion class.	13
4.6	Ground-truth and BiomedParse-predicted localization for consolidation.	13
4.7	Ground-truth and LC-KSVD-predicted localization for consolidation.	14

List of Tables

Acronyms

AC - Alternating Current
DC - Direct Current
LV - Low Voltage
MV - Medium Voltage

Chapter 1

Introduction

Developing a unified architectural pipeline for multi-class abnormality classification and grounded localization in 3D Computed Tomography (CT) addresses two of the most critical imperatives in clinical AI adoption: diagnostic accuracy and explainability. Specifically, utilizing dense pixel-level segmentation as the mechanism for localization forces the architecture to provide voxel-level visual evidence for its diagnostic claims. This ensures that the model is held accountable and its reasoning aligns with clinical diagnosis.

1.1 Background

The transition from 2D projectional radiography to volumetric CT represents a paradigm shift in medical image analysis. For Lung disease diagnosis, Chest X-rays (CXR) are commonly used as the first imaging test due to its low cost and low radiation exposure. However, studies show that their capability to detect abnormalities is limited when compared to 3D CT. [1].

Furthermore, CXR has historically benefited from the availability of massive, radiologist-annotated datasets, while the analysis of chest CT volumes has been severely constrained by a lack of such data. However, the recent introduction of large-scale, text-paired datasets has facilitated the development of state-of-the-art

(SotA) vision-language foundation models tailored specifically for volumetric imaging. [2] [3].

Beyond deep learning approaches, Dictionary Learning algorithms, particularly, Label Consistent KSVD (LC-KSVD) [4], Label Embedded Dictionary Learning (LEDL) and Fischer Discriminant Dictionary Learning (FDDL) offer viable approaches for interpretable CT analysis. All three techniques incorporate label information during the dictionary training process, and ensures that the sparse representation of the data is highly discriminative, making them well-suited for multi-abnormality tasks.

1.2 Problem Statement

The clinical utility of Deep Learning-based approaches is hindered by a lack of interpretability. While tools like Grad-CAM provide post-hoc explanations to model decisions, these are often insufficient for medical professionals who require a granular understanding of the model’s decision making process. In high stakes environments, a model that produces a diagnosis without transparency in its reasoning is difficult

to integrate into standard workflows, as it lacks the accountability necessary for legal and ethical compliance.

1.3 Objectives and Scope

The primary goal of this project is to address the limitations of existing Deep Learning based approaches for Lung abnormality classification and localization by developing an inherently interpretable dictionary learning framework. This approach aims to enable transparent, anatomically grounded decision making, where each classification outcome can be traced back to specific pathological patterns. This aligns with clinical reasoning and facilitates validation and adoption in radiological workflows. The specific objectives are as follows:

1.3.1 Objectives

- Implement a Dictionary Learning approach to learn normal and pathological dictionary atoms using a publicly available annotated 3D chest CT dataset.
- Map dictionary atoms back to anatomical features and visualize their spatial contribution maps.
- Benchmark model performance against state-of-the-art classification and segmentation models on the same dataset.
- Evaluate abnormality localization by the model on a real-world dataset with expert radiologist review.

1.3.2 Scope

The scope of this study focuses on the development and evaluation of a dictionary learning based approach for classification and grounded localization of lung abnormalities using volumetric 3D CT data, with the following limitations:

- This study is limited to 3 thoracic abnormalities: Pulmonary Nodules, Ground-Glass Opacities, and Consolidations, based on expert radiologist recommendation and dataset availability.
- This study will utilize a single publicly available dataset for training and evaluation, which may limit generalizability to other populations and scanner types.
- The proposed approach will be limited only to linear methods, and will not explore non-linear or hybrid models that provide better performance for less explainability.

1.4 Literature Review

The landscape of 3D CT analysis is currently divided into three primary architectural paradigms:

- Classification-based models that prioritize diagnostic labels

- Segmentation-based models that prioritize boundary precision
- Unified architectures designed for grounded localization

SotA models are predominantly specialized in either classification or segmentation. However, recent systems, particularly in 2025 and 2026, are moving toward multi-task learning or foundation models that can perform both diagnosis and voxel-level segmentation simultaneously.

Table 1 summarizes the key models in each category, their primary task orientation, and their performance on standard benchmarks.

1.4.1 Classification Architectures

Classification-based models for 3D CT have shifted away from 3D Convolutional Neural Networks (CNNs) such as 3D U-Net, 3D ResNet and DenseNet pipelines, toward Vision Transformers (ViTs) and self-supervised encoders. The benchmark for this paradigm is the CT-RATE dataset, which introduced the first public large-scale 3D medical imaging dataset with paired textual reports, enabling the development of Contrastive Language-Image Pretraining for Computed Tomography (CT-CLIP).

CT-CLIP is a VLFM designed for 3D chest CT analysis. It extends the CLIP framework to volumetric medical imaging by jointly learning representations from CT scans and their corresponding radiology reports. [5, 6] Through contrastive learning, the model aligns image features with textual descriptions, enabling tasks such as zero-shot classification, cross-modal retrieval, and report-guided abnormality localization. [5] The model is particularly effective in learning semantic relationships between radiological findings and volumetric imaging patterns such as GGO, consolidation, fibrosis, and pulmonary nodules. [5, 7]

Recent benchmarks such as Merlin and DINOv3 have further challenged the necessity of domain-specific pretraining. Merlin is a VLFM exclusively pre-trained on abdominal CT scans. Despite its training, Merlin has consistently demonstrated exceptional zero-shot performance on chest CT classification tasks, even when it was restricted to a linear probe, significantly outperforming models like CT-CLIP, M3FM, and CT-FM that were explicitly trained on thoracic data, by an average of 12.3% in Area Under the Receiver Operating Characteristic Curve (AUROC). Specifically, in cross-domain retrieval tasks, Merlin surpassed CT-CLIP by 24.7%.

DINOv3, a self-supervised ViT pretrained on over 1.7 billion natural images, outperform specialized medical models like CT-Net and CT-CLIP across all metrics on several classification tasks due to its ability to capture fine-grained textural features.

1.4.2 Segmentation Architectures

The prevailing SotA for supervised volumetric segmentation remains variations of the 3D U-Net, such as nnU-Net, which utilizes an adaptive framework to optimize preprocessing and architecture for specific datasets. However, the field has recently shifted toward foundation models capable of universal segmentation such as SegVol and SAM-Med3D. The SegVol model is the first 3D foundation model specifically designed for universal volumetric medical image segmentation. Similarly, SAM-Med3D adapts the Segment Anything Model (SAM) for 3D medical data.

Recent evaluations published at the Machine Learning for Healthcare (MLHC) Conference in 2025, which specifically utilized the ReXGroundingCT dataset, revealed that the current SotA text-prompted segmentation models that demonstrate near-perfect performance in general computer vision tasks struggle with the nuanced clinical descriptions and ill-defined boundaries of chest CT abnormalities.

1.4.3 Unified Architectures

Recent models have been designed to integrate abnormality detection, segmentation, and multimodal chain-of-thought reasoning, enabling pixel-level localization, structured report generation, and physician-like diagnostic inference in a single framework. BiomedParse, and its cloud-deployed version MedImageParse, developed by Microsoft, are foundation models designed to parse 3D medical images. Other notable models include Citrus-V, MTMed3D and MedVista3D, where each architecture incorporates a multi-task learning (MTL) framework.

Table 1

1.4.4 Dictionary Learning Approaches

Applications of sparse dictionary learning in medical imaging have demonstrated significant success across diverse diagnostic tasks. In histopathological analysis, Discriminative Feature oriented Dictionary Learning (DFDL) has shown the ability to learn class specific dictionaries that achieve high classification accuracy even with limited training data a critical advantage when generous training samples are not available[19]. In neuroimaging, multi feature kernel supervised dictionary learning approaches have demonstrated superior performance in distinguishing Alzheimer’s disease patients from cognitively unimpaired individuals, achieving 98.18% classification accuracy while maintaining fast testing times[20].

Most relevant to our work, dictionary learning has been successfully applied to lung disease diagnosis from CT images. One approach utilized both frozen dictionary learning and label consistent LC-KSVD to classify COVID 19 CT scans, achieving 89% accuracy on the CC-CCII dataset[21]. The framework’s key advantage lies in its interpretability, where the sparse coding output can be directly mapped back to the original input signal domain. This work demonstrates that sparse representation methods offer an effective balance between performance and explainability for lung disease classification tasks, providing clear feature representations while maintaining competitive accuracy. Similarly, sparse representation methods have achieved 95.4% accuracy in classifying diffuse lung disease patterns on HRCT images, effectively handling the high variety and complexity of patterns[22].

1.4.5 Explainability Approaches

Grad-CAM is one of the most widely used post-hoc explanation methods for CNN-based classifiers [8]. It generates class-discriminative heatmaps by using gradients flowing into the final convolutional layer, allowing the user to identify image regions that influenced a prediction [8]. Its main limitation is that it produces low-resolution saliency maps and does not directly provide precise lesion boundaries or voxel-level localization, which is a major drawback for 3D CT interpretation [8?].

Seg-Grad-CAM extends this idea by adapting Grad-CAM to semantic segmentation outputs [?]. Instead of only highlighting regions that are important for image-level classification, it produces pixel-level relevance heatmaps for segmentation predictions [?]. This is particularly useful for volumetric medical imaging tasks where the target abnormality occupies a sparse region within a large 3D scan.

HiRes-CAM addresses one of the key weaknesses of Grad-CAM by producing higher-resolution and more faithful localization maps [?]. This is important in chest CT because many abnormalities, such as small nodules, consolidations, or subtle interstitial changes, may be missed by coarse heatmaps. HiRes-CAM aims to preserve fine spatial detail, making it more suitable for medical imaging applications where precise localization is required [?]. However, like other gradient-based methods, it remains a post-hoc explanation technique and does not change the internal decision process of the model.

1.5 Report Outline

This report is organized into five main chapters.

Chapter 1 introduces the research problem, background, objectives, scope, and literature review related to multi-class abnormality classification and grounded localization in 3D chest CT imaging. It also discusses existing deep learning and dictionary learning approaches together with their limitations in explainability.

Chapter 2 presents the theoretical background required for this study, including sparse representation, dictionary learning, K-SVD, LC-KSVD, frozen dictionary learning, and Vision Transformers used in medical imaging. In addition, important radiological abnormalities such as Ground Glass Opacity (GGO), consolidation, and pulmonary nodules are discussed.

Chapter 3 describes the proposed methodology of the project. This includes dataset acquisition, preprocessing steps, patch extraction, implementation of the LC-KSVD framework, abnormality classification pipeline, and the proposed back-projection mechanism used for interpretable localization and explainability generation.

Chapter 4 explains the experimentation and evaluation framework used in this research. It presents the selected baseline models, evaluation metrics, and the three evaluation axes: classification accuracy, grounded localization quality, and explainability. The chapter also summarizes the current progress of the project and planned evaluations.

Finally, Chapter 5 discusses the overall findings, contributions, limitations, and future improvements of the proposed interpretable dictionary learning framework for 3D chest CT abnormality analysis.

Chapter 2

Background

2.1 Dictionary Learning

Dictionary learning is a signal modeling technique where a CT image patch is expressed as a linear combination of a small number of basis elements called dictionary atoms. [9] Instead of using all available features, the signal is approximated as:

$$y = Dx$$

where D represents the dictionary and x is a sparse coefficient vector containing mostly zero values. This produces a compact and interpretable representation. [10]

The dictionary D is typically learned using K-SVD (K-Singular Value Decomposition), which solves the following optimization problem:

$$\min_{D, X} \|P - DX\|_F^2 \quad \text{s.t.} \quad \forall i, \|x_i\|_0 \leq T$$

where P represents the training data, X is the sparse coefficient matrix, and T controls the sparsity level. K-SVD alternates between sparse coding and dictionary update steps to minimize reconstruction error. However, it does not incorporate class labels, which limits its discriminative ability for classification tasks. [9]

It jointly learns a dictionary, sparse codes, and a classifier so that the learned representations are not only good for reconstruction but also discriminative for classification through label consistency constraints. However, LC-KSVD still learns a single shared dictionary for all classes.

To address this limitation, we adopt a two-stage dictionary learning strategy inspired by constrained dictionary learning approaches. [11] In the first stage, a dictionary is learned from normal CT scans to capture healthy anatomical structures. In the second stage, this dictionary is frozen and kept fixed, while additional atoms are learned only from abnormal scans. This approach can be interpreted as:

- preserving D_{normal} learned from healthy data
- learning D_{abnormal} independently for disease patterns

This separation ensures that normal anatomical structures are not overwritten, while abnormal features are explicitly isolated, improving interpretability and stability.

2.2 Radiological Patterns in Lung CT Imaging

2.2.1 CT Imaging

Computed Tomography (CT) imaging is widely used in lung disease diagnosis due to its ability to generate detailed cross-sectional images of lung structures. [12] CT scans provide high-resolution visualization of abnormalities such as Ground Glass Opacity (GGO), consolidation, and pulmonary nodules, enabling accurate assessment and early detection of lung diseases. [13]

In modern medical imaging systems, CT imaging also supports computer-aided diagnosis and deep learning-based analysis for automated detection, segmentation, and classification of lung abnormalities. [14] These capabilities improve diagnostic efficiency and assist radiologists in clinical decision-making.

2.2.2 Ground Glass Opacity (GGO) and Consolidation

Ground Glass Opacity (GGO) and consolidation are two important radiological patterns used in lung CT diagnosis. [13]

Ground Glass Opacity (GGO): [15]

- Appears as hazy regions in CT images
- Does not fully obscure underlying blood vessels or airway structures
- Typically indicates partial filling of air spaces, inflammation, or early-stage infection

Consolidation: [13]

- Appears denser than GGO
- Completely obscures underlying lung structures such as vessels and airway walls
- Often associated with severe infection, fluid accumulation, or advanced disease stages

The transition between GGO and consolidation is clinically important, as it can indicate disease progression or worsening lung conditions. [16]

2.2.3 Pulmonary Nodules

Pulmonary nodules are small, rounded lesions commonly observed in lung CT scans and are important indicators for early disease detection. [17]

Characteristics of Nodules: [17]

- Small, well-defined or slightly irregular rounded opacities
- Usually less than 3 cm in diameter
- Can appear solid, subsolid, or ground-glass in nature

Clinical Significance:

- May represent benign conditions such as infections or scars [18]
- Can also indicate early-stage lung cancer in suspicious cases [19]
- Requires careful monitoring of size, shape, and growth over time [17]

Pulmonary nodules are critical in early screening and risk assessment, especially in automated CT analysis systems. [20]

2.3 Vision Transformers (ViTs)

Vision Transformers (ViTs) are deep learning architectures that adapt the transformer mechanism, originally developed for natural language processing (NLP), to image analysis tasks. [21, 22] Vision Transformers divide an image into smaller patches and process them as sequential tokens. [22]

Given an input image I , it is divided into fixed-size patches. Each patch is flattened and projected into an embedding vector. Positional embeddings are then added to preserve spatial information before the sequence is passed into transformer encoder layers.

The self-attention mechanism used in Vision Transformers allows the model to capture long-range dependencies and global contextual relationships between different regions of an image. [22] The attention operation can be represented as:

$$\text{Attention}(Q, K, V) = \text{Softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where Q represents queries, K represents keys, V represents values, and d_k is the dimension of the key vectors.

In medical imaging, Vision Transformers have shown strong performance in tasks such as classification, segmentation, abnormality detection, and disease localization. [23] Their ability to model global image context is particularly useful for analyzing complex CT volumes where abnormalities may span multiple regions. [24]

For lung CT analysis, Vision Transformer based architectures have shown strong capability in learning global representations of radiological patterns such as Ground Glass Opacity (GGO), consolidation, and pulmonary nodules. [22, 23] Hybrid architectures combining CNNs and Vision Transformers are also widely used to capture both local texture information and global contextual features. [25]

Recent developments such as Swin Transformers and 3D Vision Transformers further extend this capability by enabling hierarchical feature extraction and volumetric analysis for 3D medical imaging applications. [26]

Chapter 3

Methodology

3.1 Data Acquisition and Preprocessing

Two public chest CT datasets: CT-RATE[27] and ReXGroundingCT[7] were used in this work. CT-RATE was used as the primary source of volume-level labels (normal/abnormal and abnormality presence) derived from the dataset provided multi finding annotations, while ReXGroundingCT was used to obtain lesion level spatial supervision through radiologist linked segmentation masks together with finding descriptions[7]. In practice, CT-RATE provided the global abnormality labels needed to build classification splits, and ReXGroundingCT provided grounded masks and metadata required for localisation-oriented experiments.

After an initial review of the 18 thoracic findings available in CT-RATE, a clinical discussion with the supervising radiologist produced a shortlist of abnormalities considered both clinically meaningful and suitable for interpretable analysis (Atelectasis, Lung nodule, Ground-Glass Opacity (GGO), Consolidation, Mosaic attenuation, Bronchiectasis, and Calcification). However, after checking practical label and mask availability across the two datasets, the final experiments were restricted to three abnormalities with the most consistent coverage: *Lung nodule*, *Lung opacity / GGO*, and *Consolidation*. The remaining candidates were excluded mainly due to limited grounded lesion samples, inconsistent annotation coverage, or insufficiently distinct mask supervision for reliable training and evaluation.

The preprocessing pipeline followed the LC-KSVD implementation. CT volumes and corresponding masks were loaded from NIfTI files, then resampled to an isotropic spacing of 1.5 mm to reduce inter-scan resolution variability. Volumes were resampled with trilinear interpolation and segmentation masks with nearest-neighbour interpolation to preserve discrete labels. Intensities were windowed using a lung HU range of $[-1000, 200]$ and linearly normalised to $[0, 1]$. For dictionary learning, each pre-processed volume was converted into fixed-size 3D patches of $32 \times 32 \times 32$ voxel. Positive patches were sampled from lesion regions using the grounded masks (minimum overlap ratio 5%, with a fallback for very small lesions). Negative patches were sampled from non overlapping regions (including normal scans), and patch vectors were L_2 -normalised prior to sparse coding.

3.2 Implementation of the LC-KSVD Approach

To implement the Dictionary Learning approach for our study, we required a concrete implementation of the LC-KSVD algorithm. [4] Currently, many of the available standard implementations of KSVD and LC-KSVD are developed using MATLAB, which is not ideal for integration with our Python-based data processing and evaluation pipeline. Although a few open-source Python implementations of LC-KSVD are available, they either lacked an importable library, or did not provide the flexibility required for our study.

Therefore, we developed a custom implementation of the LC-KSVD algorithm in Python and published it on PyPI under the name “reppi.” The implementation follows a modular design, consisting of three separate components:

1. SRC - Utility Functions shared across Sparse Coding Algorithms
2. OMP - Implementation of Batch-OMP and OMP-Cholesky approaches for Sparse Coding
3. KSVD - Implementation of KSVD Dictionary Learning Algorithm
4. LC-KSVD - Implementation of LC-KSVD1 and LC-KSVD2 Dictionary Learning Algorithms

This design allows for flexibility in integrating alternative sparse coding methods. The source code is available at github.com/ckekula/reppi.

This modular architecture enables straightforward integration of the algorithm through library imports, while providing the flexibility required to experiment with different hyperparameters and configurations. The implementations of the SRC, OMP, and LC-KSVD modules are provided in Appendix A, Appendix B, and Appendix C, respectively.

3.3 Abnormality Classification

Abnormality classification was implemented as a set of one-vs-normal binary LC-KSVD2 models, trained separately for each target finding (Lung nodule, Lung opacity, and Consolidation). After preprocessing, each CT volume is represented as a matrix of vectorised $32 \times 32 \times 32$ patches, with each patch column-wise normalised to unit L_2 norm before sparse coding. At inference time, patch-wise sparse codes Γ are computed using OMP with the learned dictionary D , and classification is performed using the learned linear classifier weights W from LC-KSVD2 by scoring $s = W_{\text{pos}} \Gamma$. For volume-level prediction, the implementation aggregates evidence by averaging the sparse codes across all sampled patches and applies a sign threshold ($s > 0$) to obtain the final abnormality present/absent decision.

3.4 Back-Projecting Class-Discriminative Atoms

The unique intuition we propose in our approach is to generate spatial contribution maps through back-projection. Because each atom in the dictionary D is associated with a specific class label through the matrix Q [4], the model can theoretically

”explain” its decision for a test patch by examining which atoms were used in its sparse representation.

To generate a spatial contribution map for a class c , the algorithm identifies all atoms d_k that have been assigned to that class. For a query voxel or patch with sparse coefficients x , the contribution C_c is calculated as the partial reconstruction:

$$C_c = \sum_{k \in \text{Class } c} d_k x_k$$

This process should theoretically back-project the class-specific features into the original voxel space, creating a map that highlights the regions contributing to the detection of a specific abnormality. for class-discriminative atoms operating on raw voxel-space inputs, the back-projection is an exact, closed-form spatial attribution.

3.5 Deriving Explainability through Segmentation Masks

To evaluate the clinical accuracy of LC-KSVD, these contribution maps are converted into binary segmentation masks for comparison with ground truth annotations. This conversion can be achieved through various thresholding methods.

Grad-CAM’s approximation has two sources of imprecision — the gradient itself is a local linear approximation of a non-linear function, and the spatial upsampling from the coarse feature map to input resolution introduces further interpolation error. LC-KSVD’s contribution map has only one source of imprecision: the thresholding step you apply to binarize it into a mask. The map itself — before thresholding — is exact given the sparse code and the dictionary.

Unlike the passive Grad-CAM approach, atom back-projection is generative. The atoms are what the model actually used to reconstruct and classify the signal. There’s no proxy involved. The localization evidence is produced by the same computation that produced the classification, not by a separate post-hoc analysis run afterwards.

Chapter 4

Evaluation Framework and Results

4.1 Evaluation Framework Design

Comparing a sparse linear approach with state-of-the-art non-linear model architectures using only raw performance scores can give an incomplete view of model behavior. In medical imaging, a useful model should not only predict the correct abnormality, but also identify the relevant anatomical region and provide a meaningful explanation for its decision. Therefore, the evaluation framework in this study is organized around three complementary evaluation axes:

- **Axis A: Classification Accuracy** evaluates the ability of the model to detect the presence or absence of each abnormality using standard discriminative metrics such as AUROC, precision, recall, and F1 score.
- **Axis B: Grounded Localization Quality** evaluates whether the model localizes the abnormality in the correct region of the 3D CT volume using spatial overlap and hit-based localization metrics such as DSC, IoU, and HIT@K.
- **Axis C: Interpretability and Explanation Faithfulness** evaluates whether the model’s explanation or attribution map is spatially aligned with the true abnormality region using metrics such as Attribution-Mask IoU, Pointing Game Accuracy, and Energy Inside Mask.

Axis C is particularly important for this study because the proposed LC-KSVD-based approach is designed to provide interpretable decision-making through dictionary atoms and sparse activations. In contrast, most deep learning and foundation models require post-hoc attribution techniques such as Grad-CAM, HiResCAM, or saliency-based methods to explain their predictions. These post-hoc methods provide useful approximations, but they do not directly represent the internal decision process of the model. Therefore, the evaluation framework explicitly separates predictive performance, localization quality, and explanation faithfulness to better highlight the strengths and limitations of each model type.

4.2 Evaluation Metrics

A common set of quantitative metrics was used to evaluate the models across the three axes. This ensures that different models are compared under the same evaluation




Figure 4.1: Architecture diagram of the proposed evaluation framework.

criteria, even though their raw outputs may differ. Classification-oriented models are evaluated using disease probability scores, localization-oriented models are evaluated using predicted masks, and interpretable models are additionally evaluated using attribution or activation maps.

4.2.1 Classification Metrics

For Axis A, the classification task measures whether the model correctly detects each abnormality. The following metrics are used:

- **Area Under the Receiver Operating Characteristic Curve (AUROC):** AUROC measures how well the model separates positive and negative cases across different decision thresholds. It is useful for imbalanced medical datasets because it evaluates ranking performance rather than relying on a single fixed threshold.
- **Precision and Recall:** Precision measures how many predicted positive cases are actually positive, while recall measures how many true positive cases are correctly detected by the model.
- **F1 Score:** F1 score is the harmonic mean of precision and recall. It provides a balanced measure when both false positives and false negatives are important.
- **Average Precision:** Average precision summarizes the precision-recall curve and is useful when positive cases are relatively rare.

4.2.2 Localization and Grounding Metrics

For Axis B, the localization task measures whether the predicted abnormality region overlaps with the ground-truth lesion mask. This is important because a model may

correctly predict the presence of a disease while still focusing on the wrong anatomical area.

- **Dice Similarity Coefficient (DSC):** Dice measures the spatial overlap between the predicted mask and the ground-truth mask. It is defined as:

$$DSC = \frac{2|X \cap Y|}{|X| + |Y|}$$

where X is the predicted mask and Y is the ground-truth mask.

- **Intersection over Union (IoU):** IoU, also known as the Jaccard index, measures the ratio between the overlapping region and the total combined region:

$$IoU = \frac{|X \cap Y|}{|X \cup Y|}$$

- **HIT@K:** HIT@K measures the proportion of cases where the model achieves at least a small amount of meaningful overlap with the ground-truth lesion region. In this work, HIT@5, HIT@10, and HIT@25 are used to indicate whether the Dice score exceeds 0.05, 0.10, and 0.25 respectively. This is useful for diffuse abnormalities where exact boundary agreement is difficult.

4.2.3 Interpretability Metrics

For Axis C, the interpretability task evaluates whether the model’s explanation is focused on the clinically relevant abnormality region. This is different from localization because the goal is not only to check where the model predicts a lesion, but also to determine whether the model’s decision is supported by the correct anatomical evidence.

- **Attribution-Mask IoU:** This measures the overlap between the thresholded attribution map and the ground-truth lesion mask. A higher value indicates that the explanation is spatially aligned with the true abnormality.
- **Pointing Game Accuracy:** This checks whether the highest-attribution point lies inside the ground-truth lesion region. If the maximum response is inside the lesion mask, the case is counted as a hit.
- **Energy Inside Mask:** This measures the percentage of total attribution energy that lies inside the ground-truth mask. A high value indicates that the model’s explanation is concentrated on the true abnormality region rather than irrelevant background structures.

4.3 Classification Accuracy Evaluation

The classification performance was evaluated across four abnormality classes: Lung Nodule, Lung Opacity, Consolidation, and Atelectasis. The proposed LC-KSVD-based baseline was compared with three foundation model baselines: Merlin, CT-CLIP, and BiomedParse. For LC-KSVD, consolidation-related performance was also used as a proxy reference for atelectasis due to current dataset and label constraints.

4.3.1 Overall Classification Performance

Figure 4.2 shows the macro-level classification performance across the evaluated models. Merlin achieved the strongest discriminative performance, with a macro-AUROC of 0.966. In contrast, CT-CLIP and BiomedParse produced lower macro-AUROC values of 0.538 and 0.509 respectively on the ReXGroundingCT evaluation setting.




Figure 4.2: Overall classification performance comparison across models.

These results suggest that Merlin is the strongest classification baseline in this setup. CT-CLIP and BiomedParse show weaker classification performance, which may be due to differences in training objectives, dataset alignment, and the use of zero-shot or proxy classification outputs. BiomedParse is mainly designed for text-guided segmentation and grounding, so its classification score should be interpreted as a proxy rather than as the output of a dedicated classifier.

4.3.2 Per-Class Analysis




Figure 4.3: Per-class AUROC comparison Across Models.

The per-class AUROC comparison in Figure 4.3 provides a more detailed view of model behavior. Merlin maintains strong performance across all evaluated classes.

The LC-KSVD baseline shows promising results for non-focal or broader abnormalities, achieving AUROC values of 0.917 for Lung Opacity and 0.806 for Consolidation. It also achieves 0.806 for Atelectasis in the current evaluation setting. However, its performance for Lung Nodule is weaker, with an AUROC of 0.524.

This behavior suggests that LC-KSVD may capture broader abnormal tissue patterns more effectively than small focal lesions. Lung nodules are compact and localized, so their detection may require finer spatial resolution, multi-scale patch modeling, or more nodule-specific training samples. In contrast, abnormalities such as lung opacity and consolidation may produce more distributed visual patterns that are more compatible with the sparse dictionary representation.

4.4 Grounded Localization Evaluation

Axis B evaluates whether the models can identify the correct abnormal region in the volumetric CT scan. This is different from classification because a model can correctly detect the presence of a disease while still localizing the wrong anatomical region. The localization evaluation therefore uses Dice, IoU, and HIT@K metrics to measure spatial agreement between predicted lesion regions and ground-truth annotations.




Figure 4.4: Overall localization performance comparison for BiomedParse and Merlin.

Figure 4.4 shows that grounded localization remains challenging for the evaluated models. The Dice, IoU, and HIT@K values indicate limited agreement between predicted lesion regions and reference annotations. This highlights an important gap between disease recognition and true spatial grounding. For clinical support, it is not sufficient to detect an abnormality; the model should also identify the correct location and extent of the abnormal region.

4.4.1 Per-Class Localization Performance

Figure 4.5 shows the per-class Dice comparison. The results indicate that localization performance varies substantially across lesion categories. This variation is likely influenced by lesion morphology, size, contrast, and boundary clarity.

Focal lesions such as lung nodules show comparatively better localization than diffuse abnormalities such as lung opacity, consolidation, and atelectasis. This suggests




Figure 4.5: Mean Dice comparison by lesion class.

that object-like lesions with clearer boundaries are easier for text-guided segmentation models to localize. In contrast, diffuse abnormalities often have ambiguous borders and variable appearance, making exact spatial overlap more difficult.

4.4.2 Qualitative Comparison of Localization Outcomes

The qualitative BiomedParse example in Figure 4.6 further illustrates the gap between abnormality detection and precise grounding. In the consolidation example, the predicted region covers a broad area of the lung, while the ground-truth annotation is more localized. This indicates that the model responds to the general lung region associated with the abnormality, but it does not accurately preserve the lesion boundary, extent, or class-specific spatial location.




Figure 4.6: Ground-truth and BiomedParse-predicted localization for consolidation.

The LC-KSVD visualization in Figure 4.7 shows a related limitation. The prediction map highlights a suspicious region in the lung, indicating some sensitivity to abnormal tissue patterns. However, the activated region is not fully aligned with the ground-truth lesion location. This suggests that the sparse representation cap-

tures some abnormal visual structure, but further refinement is needed to improve anatomical alignment.




Figure 4.7: Ground-truth and LC-KSVD-predicted localization for consolidation.

Overall, the localization results show that both foundation models and the proposed sparse representation approach face challenges in producing precise lesion-level grounding in 3D CT. This supports the need for evaluation beyond classification scores alone.

4.5 Inherent Explainability Evaluation

Axis C focuses on whether model predictions are supported by clinically meaningful evidence. This is a key motivation for the proposed LC-KSVD approach. Deep learning models such as Merlin, CT-CLIP, and BiomedParse can be evaluated with post-hoc explanation methods or grounding maps, but these explanations are not always part of the model’s internal decision process. LC-KSVD, in contrast, provides a more direct form of traceability through its dictionary atoms and sparse coefficients.

4.5.1 Atom-to-Voxel Traceability

LC-KSVD represents input patches using a learned dictionary and sparse activation coefficients. Since the reconstruction and classification processes are based on linear combinations of dictionary atoms, the contribution of selected atoms can be traced back to the corresponding voxel regions in the CT volume. This provides a direct connection between the model’s prediction and the spatial image patterns that contributed to it.

This atom-to-voxel traceability is useful because it allows the model’s decision to be inspected in terms of activated dictionary components. Instead of relying only on a final disease score, the user can examine which atoms were activated and where their contributions appear in the volume.

4.5.2 Spatial Faithfulness of Explanations

The explanation quality is evaluated using attribution maps generated from model activations or atom contributions. These maps are compared with the ground-truth lesion masks using the following metrics:

- **Attribution-Mask IoU:** Measures how much the explanation map overlaps with the true lesion region.
- **Pointing Game Accuracy:** Checks whether the strongest explanation point falls inside the annotated lesion region.
- **Energy Inside Mask:** Measures how much of the explanation signal is concentrated inside the ground-truth mask.

Together, these metrics assess whether the explanation is spatially faithful to the disease region. A model with good interpretability should not only predict the correct abnormality, but should also base its decision on the relevant anatomical evidence.

For LC-KSVD, these metrics are particularly meaningful because the explanation comes from the model structure itself through sparse codes and dictionary atoms. Even when its classification performance is lower than that of a large foundation model, its explanations can provide additional value by making the decision process more transparent and easier to inspect.

4.6 Summary of Evaluation Findings

The evaluation framework provides a multi-dimensional comparison of model performance. Axis A shows that Merlin is the strongest classification baseline, while LC-KSVD demonstrates promising performance for lung opacity, consolidation, and atelectasis. Axis B shows that localization remains difficult for all evaluated models, especially for diffuse abnormalities. BiomedParse performs relatively better on focal lung nodules, but its localization performance is inconsistent across classes. Axis C highlights the main advantage of LC-KSVD: its ability to provide interpretable atom-based explanations that can be traced back to voxel-level regions.

Overall, the results indicate that classification accuracy alone is not sufficient for evaluating medical AI models in 3D chest CT analysis. A reliable model should be accurate, spatially grounded, and explainable. The proposed evaluation framework captures these three requirements and provides a fair basis for comparing sparse interpretable models with modern foundation model baselines.

Chapter 5

Discussion

Our project investigated an interpretable sparse dictionary learning framework for abnormality classification and localization in 3D chest CT scans. A key contribution is the frozen dictionary approach, where normal anatomical patterns are learned first and abnormalities are learned separately. This helps isolate pathological features and improves interpretability.

Unlike many state-of-the-art deep learning models that rely on post-hoc explainability methods such as Grad-CAM, the proposed LC-KSVD framework provides inherent explainability. The generated predictions can be directly linked to dictionary atoms and reconstructed voxel-space regions through atom back-projection, allowing more transparent clinical interpretation.

The experimental observations show that the framework can successfully classify and localize pulmonary nodules, Ground Glass Opacity (GGO), and consolidation patterns from 3D chest CT scans while maintaining competitive performance. The proposed objectives of implementing an interpretable dictionary learning pipeline, generating localization maps, and evaluating the framework against existing approaches were successfully achieved.

In real-world clinical applications, this approach could support radiologists by providing both diagnostic predictions and anatomically grounded explanations, improving trust and transparency in AI-assisted diagnosis systems.

However, the proposed method has several limitations. The framework is currently limited to three abnormalities and does not yet provide a fully unified end-to-end classification and segmentation architecture. In addition, although the explainability is stronger, the classification accuracy may still be lower than large-scale deep learning foundation models. Future work will focus on improving performance, expanding abnormality coverage, and integrating more advanced hybrid architectures.

Appendix A

Appendix 1

The Appendices are numbered as A, B,.. and so on. Sections and Subsections of the appendix then would be A.1, A.1.1, A.2 etc.

Appendix B

Appendix 2

Bibliography

- [1] W. H. Self, D. M. Courtney, C. D. McNaughton, R. G. Wunderink, and J. A. Kline, “High discordance of chest x-ray and computed tomography for detection of pulmonary opacities in ed patients: implications for diagnosing pneumonia,” *The American Journal of Emergency Medicine*, vol. 31, no. 2, pp. 401–405, 2013.
- [2] M. Baharoon, L. Luo, M. Moritz, A. Kumar, S. E. Kim, X. Zhang, M. Zhu, M. H. Alabbad, M. S. Alhazmi, N. P. Mistry, L. Bijnens, K. R. Kleinschmidt, B. Chrisler, S. Suryadevara, S. S. D. Jaliparthi, N. M. Prudlo, M. D. Marino, J. Palacio, R. Akula, D. Zhou, H.-Y. Zhou, I. E. Hamamci, S. J. Adams, H. R. AlOmaish, and P. Rajpurkar, “Rexgroundingct: A 3d chest ct dataset for segmentation of findings from free-text reports,” 2025.
- [3] J. Chen, Y. Liang, and Z. Liu, “Dflash: Block diffusion for flash speculative decoding,” 2026.
- [4] Z. Jiang, Z. Lin, and L. S. Davis, “Label consistent k-svd: Learning a discriminative dictionary for recognition,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2013.
- [5] I. E. Hamamci, S. Er, E. Simsar, H. Yang, F. Xu, *et al.*, “Ct-clip: Learning transferable visual models from natural language supervision for ct,” *Nature Machine Intelligence*, 2024.
- [6] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, *et al.*, “Learning transferable visual models from natural language supervision,” in *International Conference on Machine Learning*, 2021.
- [7] M. Baharoon, L. Luo, M. Moritz, A. Kumar, S. E. Kim, X. Zhang, M. Zhu, M. H. Alabbad, M. S. Alhazmi, N. P. Mistry, L. Bijnens, K. R. Kleinschmidt, B. Chrisler, S. Suryadevara, S. S. D. Jaliparthi, N. M. Prudlo, M. D. Marino, J. Palacio, R. Akula, D. Zhou, H.-Y. Zhou, I. E. Hamamci, S. J. Adams, H. R. AlOmaish, and P. Rajpurkar, “Rexgroundingct: A 3d chest ct dataset for segmentation of findings from free-text reports,” 2025.
- [8] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, “Grad-cam: Visual explanations from deep networks via gradient-based localization,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pp. 618–626, 2017.
- [9] M. Aharon, M. Elad, and A. Bruckstein, “K-svd: An algorithm for designing overcomplete dictionaries for sparse representation,” *IEEE Transactions on Signal Processing*, vol. 54, no. 11, pp. 4311–4322, 2006.

- [10] M. Elad, *Sparse and Redundant Representations: From Theory to Applications in Signal and Image Processing*. Springer, 2010.
- [11] Y. Yankelevsky and M. Elad, “Dual graph regularized dictionary learning,” *IEEE Transactions on Signal Processing*, vol. 64, no. 20, pp. 5400–5415, 2016.
- [12] I. Sluimer, M. Prokop, and B. van Ginneken, “Computer analysis of ct scans of the lungs,” *IEEE Transactions on Medical Imaging*, vol. 25, no. 4, pp. 385–405, 2006.
- [13] D. M. Hansell, A. A. Bankier, H. MacMahon, T. C. McLoud, N. L. Muller, and J. Remy, “Fleischner society: Glossary of terms for thoracic imaging,” *Radiology*, vol. 246, no. 3, pp. 697–722, 2008.
- [14] G. Litjens, T. Kooi, B. E. Bejnordi, A. A. A. Setio, F. Ciompi, M. Ghafoorian, J. A. W. M. van der Laak, B. van Ginneken, and C. I. Sanchez, “A survey on deep learning in medical image analysis,” *Medical Image Analysis*, vol. 42, pp. 60–88, 2017.
- [15] H. J. Park, K. S. Lee, J. B. Hur, Y. Kim, S. H. Choi, and O. Y. Kwon, “Ground-glass opacity on high-resolution ct: A sign of active pulmonary disease,” *Korean Journal of Radiology*, vol. 12, no. 4, pp. 458–471, 2011.
- [16] M. Chung, A. Bernheim, X. Mei, X. Zhang, M. Huang, X. Zeng, Y. Yang, Z. A. Fayad, A. Jacobi, K. Li, S. Li, and H. Shan, “Ct imaging features of 2019 novel coronavirus (2019-ncov),” *Radiology*, vol. 295, no. 1, pp. 202–207, 2020.
- [17] H. MacMahon, D. P. Naidich, J. M. Goo, K. S. Lee, A. N. C. Leung, J. R. Mayo, P. Meiner, Y. Ohno, C. A. Powell, M. Prokop, G. D. Rubin, C. M. Schaefer-Prokop, W. D. Travis, P. E. van Schil, and A. A. Bankier, “Guidelines for management of incidental pulmonary nodules detected on ct images: From the fleischner society 2017,” *Radiology*, vol. 284, no. 1, pp. 228–243, 2017.
- [18] M. K. Gould, J. Donington, T. J. Lynch, P. Mazzone, and P. B. Pennell, “Evaluation of individuals with pulmonary nodules: when is it lung cancer?,” *Chest*, vol. 132, pp. 108S–130S, 2007.
- [19] D. R. Aberle, A. M. Adams, C. D. Berg, W. C. Black, J. D. Clapp, and R. M. Fagerstrom, “Reduced lung-cancer mortality with low-dose computed tomographic screening,” *New England Journal of Medicine*, vol. 365, pp. 395–409, 2011.
- [20] A. A. A. Setio, A. Traverso, T. de Bel, F. Ciompi, S. J. van Riel, E. T. Scholten, B. van Ginneken, and C. van den Bogaard, “Validation, comparison, and combination of algorithms for automatic detection of pulmonary nodules in computed tomography images: The luna16 challenge,” *Medical Image Analysis*, vol. 42, pp. S1–S13, 2017.
- [21] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems*, vol. 30, pp. 5998–6008, 2017.

- [22] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, “An image is worth 16x16 words: Transformers for image recognition at scale,” in *International Conference on Learning Representations*, 2021.
- [23] F. Shamshad, S. Khan, S. W. Zamir, M. H. Khan, M. Hayat, F. S. Khan, and H. Fu, “Transformers in medical imaging: A survey,” *Medical Image Analysis*, vol. 88, p. 102802, Aug 2023.
- [24] A. Khan, Z. Rauf, A. R. Khan, S. Rathore, S. H. Khan, N. S. Shah, U. Farooq, H. Asif, A. Asif, U. Zahoor, R. U. Khalil, S. Qamar, U. H. Asif, F. B. Khan, A. Majid, and J. Gwak, “A recent survey of vision transformers for medical image segmentation,” 2023.
- [25] J. Chen, Y. Lu, Q. Yu, X. Luo, E. Adeli, Y. Wang, L. Lu, A. L. Yuille, and Y. Zhou, “Transunet: Transformers make strong encoders for medical image segmentation,” 2021.
- [26] Z. Liu, Y. Lin, Y. Cao, H. Hu, Y. Wei, Z. Zhang, S. Lin, and B. Guo, “Swin transformer: Hierarchical vision transformer using shifted windows,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 9992–10002, 2021.
- [27] I. E. Hamamci, S. Er, C. Wang, F. Almas, A. G. Simsek, S. N. Esirgun, I. Dogan, O. F. Durugol, B. Hou, S. Shit, W. Dai, M. Xu, H. Reynaud, M. F. Dasdelen, B. Wittmann, T. Amiranashvili, E. Simsar, M. Simsar, E. B. Erdemir, A. Alanbay, A. Sekuboyina, B. Lafci, A. Kaplan, Z. Lu, M. Polacin, B. Kainz, C. Bluethgen, K. Batmanghelich, M. K. Ozdemir, and B. Menze, “Developing generalist foundation models from a multimodal dataset for 3d computed tomography,” 2025.