



Interpretable Sparse Dictionary Learning for Diagnostic Hypotheses from 3D Chest CT Findings

An undergraduate project proposal report submitted to the

Department of Electrical and Information Engineering
Faculty of Engineering
University of Ruhuna
Sri Lanka

in partial fulfillment of the requirements for the

**Degree of the Bachelor of the Science of Engineering
Honours**

by

K.R.C.S. Kekulawala	-	EG/2021/4609
M.D.Y.B. Kularathne	-	EG/2021/4620
B.G.P.U. Thilakshana	-	EG/2021/4831
B.P.G.L. Thejana	-	EG/2021/4828

.....
Dr. Kaveen Liyanage
(Supervisor)

.....
Mrs. Sithara Mahagama
(Co-Supervisor)

Abstract

Lung disease diagnosis from 3D chest CT scans faces critical challenges in computational efficiency and clinical interpretability. Traditional 3D Convolutional Neural Networks require extensive labeled data and GPU resources, while post-hoc explainability methods like Grad-CAM provide insufficient granularity for medical decision making. While deep learning models achieve high accuracy, their black box nature and substantial computational demands limit clinical adoption. This project proposes a sparse dictionary learning framework that addresses these limitations through inherently interpretable feature representations. The methodology uses frozen dictionary learning approach combined with Label-Consistent K-SVD (LC-KSVD) to capture both normal anatomical patterns and disease-specific pathological features. Dictionary atoms are mapped to clinically meaningful structures such as Ground-Glass Opacity (GGO) and consolidation, enabling transparent diagnostic reasoning. Each dictionary atom corresponds to specific volumetric patterns, allowing clinicians to trace classification decisions back to anatomical features rather than relying on post-hoc visualization techniques. The proposed system will be benchmarked against standard K-SVD and 3D CNN baselines using the CT-RATE and ReXGroundingCT datasets. The sparse representation framework naturally handles class imbalance and requires fewer training samples than deep learning approaches, making it particularly suited for rare pathologies. This research aims to demonstrate that sparse representation methods can provide clinically interpretable lung disease classification while maintaining competitive performance with reduced computational overhead.

Contents

Acronyms	vi
1 Introduction	1
1.1 Background	1
1.2 Problem Statement	3
1.2.1 Computational Complexity	3
1.2.2 Interpretability Bottleneck	3
1.3 Objectives and Scope	3
1.3.1 Objectives	3
1.3.2 Scope	4
2 Literature Review	5
2.1 Previous Work	5
2.1.1 Deep Learning for Lung Disease Diagnosis	5
2.1.2 Sparse Dictionary Learning in Medical Imaging	5
2.1.3 Explainable AI in Medical Imaging	7
2.1.4 Volumetric Clinical Feature Extraction	8
2.2 Gaps in Literature	8
2.2.1 Inherently Interpretable Architectures	8
2.2.2 Frozen Dictionary Learning for Pathological Differentiation	9
2.2.3 Handling Class Imbalance with Interpretable Models	10
3 Methodology	11
3.1 Data acquisition and pre-processing	11
3.2 Model Training	12
3.3 Classification and Disease Diagnosis	12
3.4 Interpretability Analysis and Visualization	13
3.5 Volumetric Feature Extraction and Quantification	13
3.6 Baseline Comparisons and Benchmarking	13
3.7 Evaluation Metrics and Validation Strategy	14
4 Timeline and Resources Required	15
4.1 Timeline	15
4.2 Resources Required	15
5 Conclusion	18
References	19

List of Figures

1.1	Examples of GGO and consolidation	2
1.2	3D CT Imaging and visualization process of the craniofacial region.	2
2.1	Visualization of Sparse Representation	6
3.1	Overall Methodology	11
3.2	Model Training Workflow	12

List of Tables

2.1	Gaps addressed by proposed approach vs existing solutions	8
3.1	Characteristics of CT-RATE and RexGroundingCT datasets	11
4.1	Overall Project Time Plan	16
4.2	Resources Required for the Project	17

Acronyms

AI Artificial Intelligence

C/G Consolidation-to-Ground-Glass-Opacity

CNN Convolutional Neural Network

CT Computed Tomography

GGO Ground-Glass Opacity

HU Hounsfield Unit

KSVD K-Singular Value Decomposition

LC-KSVD Label-Consistent K-Singular Value Decomposition

XAI Explainable Artificial Intelligence

Chapter 1

Introduction

1.1 Background

Lung disease remains a leading cause of global morbidity and mortality[1]. In the diagnosis of lung disease, Chest X-rays (CXR) are commonly used as the first imaging test due to its low cost and low radiation exposure. However, their diagnostic sensitivity is limited by their two-dimensional nature, which causes anatomical superposition, allowing visualization of only about 70% of the lung volume[2]. Computed Tomography (CT) is a radiographic technology that combines X-ray technology with computer reconstruction to produce images of cross sections of the human body[3]. This overcomes many of the limitations of radiography by eliminating anatomical superposition and enabling full volumetric assessment of the lungs.

CT produces high-resolution slices that reveal fine structural details such as normal parenchyma, Ground-Glass Opacity (GGO), and consolidation (figure 1.1). The traditional approach for CT analysis relies on a slice-by-slice approach where a radiologist views a series of axial cross-sections and identifies the index slice (the single image where a pathology appears at its largest diameter)[4]. A 2D slice only captures a small fraction of a 3D structure. Studies show that 2D measurements are highly sensitive to the patient’s orientation and the specific slice chosen, leading to high inter-observer variability[4]. 2D models also treat each slice as an independent data point, losing the spatial continuity between adjacent tissues.

3D volumetric analysis represents a quantitative shift, wherein the entire dataset is used to reconstruct the lung’s anatomy as a solid volume (figure 1.2). This enables more accurate characterization of lesion size, shape, and spatial distribution, while reducing sensitivity to slice selection and viewing angle. 3D representations better capture disease morphology and progression, providing a more robust visualization for quantitative analysis.

Modern radiology increasingly relies on Artificial Intelligence (AI) to streamline the acquisition of radiological analysis on chest CTs[6]. However, the exponential increase in data volume generated by 3D CT scanners presents a significant challenge for human interpretation and computational analysis. The prevailing research trend has focused heavily on the deployment of 3D deep learning architectures, such as Convolutional Neural Networks (CNNs), achieving remarkable success in tasks such as segmenting suspected infection regions and differentiat-




Figure 1.1: Examples of GGO and consolidation on axial slices of lung CT scans from COVID-19 patients[5]. Arrows point to the areas affected by the infection. Left: GGO Right: Consolidation.




Figure 1.2: 3D CT Imaging and visualization process of the craniofacial region.

ing between various types of carcinoma[7]. Despite these advancements, these models frequently encounter challenges regarding computational complexity, data scarcity, and a lack of interpretability. The research proposed herein advocates for a departure from these black-box paradigms, suggesting that sparse representation methods offer a more intuitive framework for the volumetric identification of lung diseases by ensuring clinical interpretability.

1.2 Problem Statement

1.2.1 Computational Complexity

Processing a 3D CT scan involves treating stacked 2D slices as a single volumetric matrix, which demands significantly more GPU memory and processing power than standard image analysis. Unlike ordinary RGB images, CT scans consist of single-channel data with varying voxel dimensions and resolution, making spatial alignment and generalizability across different scanners more difficult. Furthermore, the curse of dimensionality remains an issue; as the network depth increases to capture global features, the risk of gradient explosion and slow convergence becomes significant.

1.2.2 Interpretability Bottleneck

Beyond computational limits, the clinical utility of deep learning is hindered by a lack of interpretability. While tools like Grad-CAM[8] provide visual heatmaps to localize lesions, these post-hoc explanations are often insufficient for medical professionals who require a granular understanding of the model’s decision-making process. In high-stakes environments, a model that produces a diagnosis without transparency in its reasoning is difficult to integrate into standard workflows, as it lacks the accountability necessary for legal and ethical compliance. This transparency gap has led to a renewed interest in models that are interpretable by design, where the internal logic is clear from the beginning.

1.3 Objectives and Scope

1.3.1 Objectives

The primary goal of this project is to address the limitations of existing AI-based lung disease diagnosis systems, specifically high computational complexity and lack of clinical interpretability, by developing an inherently interpretable and efficient dictionary learning framework. The proposed approach aims to enable transparent, anatomically-grounded decision making, where each classification outcome can be traced back to specific pathological patterns. This aligns with clinical reasoning and facilitates validation and adoption in radiological workflows. The specific objectives are as follows:

1. Develop an interpretable dictionary learning framework for lung disease classification from 3D chest CT scans
 - Implement a training strategy to learn normal and pathological dictionary atoms using a publicly available annotated 3D chest CT dataset
 - Integrate label consistency to improve class discrimination
2. Improve model interpretability through visualization of learned dictionary atoms

- Visualize dictionary atoms in 3D and map them to anatomical and pathological lung features
 - Analyze and compare sparse coefficient patterns across different disease classes
3. Evaluate the effectiveness of the proposed framework
 - Benchmark performance against standard KSVD without label consistency
 - Compare results with a 3D CNN baseline (3D-ResNet)
 4. Integrate clinically relevant volumetric features into the analysis
 - Perform segmentation to identify radiological features such as GGO and consolidation
 - Compute the Consolidation-to-GGO (C/G) ratio and relate model predictions to clinical thresholds
 5. Optionally include verification from radiologists to support the analysis
 - Gather opinions from multiple radiologists on the regions identified by the model
 - Utilize this feedback to understand how well the model’s predictions align with real clinical judgment

1.3.2 Scope

The scope of this study focuses on the development and evaluation of a dictionary learning-based framework for diagnostic hypotheses using volumetric 3D chest CT data, with the following limitations:

- The proposed solution is limited to volumetric CT scans of the thorax and does not consider other imaging modalities such as chest X-rays, MRI, or PET scans.
- The methodology is restricted to dictionary learning-based approaches, with particular emphasis on KSVD and LC-KSVD. Deep learning models such as CNNs are not the primary focus of this work and are included only as baseline comparators for performance benchmarking.
- Clinically, the analysis is confined to the diagnostic hypotheses of 14 disease categories, based on the 18 abnormalities annotated in the CT-RATE dataset. This list does not include all possible differentials but covers the most clinically relevant categories for an AI diagnostic model. The study does not address disease prognosis or treatment outcome prediction.
- The scope excludes real-time clinical deployment, cross-institutional generalization studies, and regulatory validation. The findings are intended to demonstrate the feasibility and interpretability of sparse representation learning for 3D lung disease analysis rather than to provide a clinically deployable diagnostic system.

Chapter 2

Literature Review

2.1 Previous Work

2.1.1 Deep Learning for Lung Disease Diagnosis

Numerous techniques have emerged in the field of automated lung disease diagnosis, each attempting to balance accuracy with computational feasibility. Early approaches relied on manual feature engineering and classical classifiers like Support Vector Machines (SVM) and Random Forests.[9] While these methods provided strong interpretability, they often lacked the flexibility to handle the noise and variance found in CT data[10]. The emergence of deep learning fundamentally changed the landscape of CT analysis, especially in areas such as[11];

- Image Segmentation and Classification
- Advancing Diagnostics with AI and CAD Systems
- Prognostics with Radiomics and Predictive Analytics
- Workflow Optimisation Using AI

Convolutional Neural Networks (CNNs) have emerged as the dominant paradigm for automated analysis of CT scans[12]. Recent surveys indicate that deep learning models achieve classification accuracy exceeding 98% for multi-disease detection tasks[13], representing a significant improvement over traditional machine learning approaches. Contemporary research has increasingly focused on 3D CNN architectures that treat CT scans as complete volumetric data rather than independent 2D slices. This volumetric approach captures spatial relationships within lung tissue more effectively, leading to improved sensitivity in detecting subtle pathological changes. Li et al. (2024)[14] demonstrated that preprocessing techniques such as lung region extraction, combined with advanced backbone architectures like ResNeSt50, significantly enhance classification performance for multiple lung pathologies [15].

2.1.2 Sparse Dictionary Learning in Medical Imaging

Sparse dictionary learning represents an unsupervised machine learning paradigm for extracting interpretable features from medical images. This technique learns

a set of dictionary elements or basis vectors that represent data in a sparse manner. In the context of 3D CT, a volumetric patch $y \in \mathbb{R}$ can be approximated by $y \sim Dx$, where $D \in \mathbb{R}^{n \times K}$ is the dictionary and $x \in \mathbb{R}^k$ is a sparse vector containing only a small number of non-zero coefficients (figure 2.1). The strength of dictionary learning lies in its ability to extract local, structural information from noisy medical images while requiring less labeled data than purely supervised deep learning approaches. This characteristic is particularly valuable for rare diseases with limited training samples[16].




Figure 2.1: Brief overview of the sparse representation[17]. (a) Sparse linear combination Dx approximating y , where $\|x\|_0 = 6$ and (b) two matrices D and X calculated in K-SVD algorithm 36 from Y , where D is a dictionary matrix, X is a matrix including representation coefficient vectors, and Y is a matrix including target signals.

The success of these models hinges on the quality of the learned dictionary. Traditional methods like KSVD focus primarily on minimizing the reconstruction error, which is effective for image denoising and super-resolution but lacks the discriminative power required for classification[18]. To address this, supervised dictionary learning techniques have emerged, which incorporate label information directly into the dictionary construction process. The LC-KSVD algorithm enhances the traditional KSVD by adding a label consistency constraint and a classification error term to the objective function[18]. This approach forces im-

ages from the same class to have similar sparse codes, thereby increasing the discriminative power of the model.

Applications of sparse dictionary learning in medical imaging have demonstrated significant success across diverse diagnostic tasks. In histopathological analysis, Discriminative Feature-oriented Dictionary Learning (DFDL) has shown the ability to learn class-specific dictionaries that achieve high classification accuracy even with limited training data—a critical advantage when generous training samples are not available[19]. In neuroimaging, multi-feature kernel supervised dictionary learning approaches have demonstrated superior performance in distinguishing Alzheimer’s disease patients from cognitively unimpaired individuals, achieving 98.18% classification accuracy while maintaining fast testing times[20].

Most relevant to our work, dictionary learning has been successfully applied to lung disease diagnosis from CT images. One approach utilized both frozen dictionary learning and label-consistent LC-KSVD to classify COVID-19 CT scans, achieving 89% accuracy on the CC-CCII dataset[21]. The framework’s key advantage lies in its interpretability, where the sparse coding output can be directly mapped back to the original input signal domain. This work demonstrates that sparse representation methods offer an effective balance between performance and explainability for lung disease classification tasks, providing clear feature representations while maintaining competitive accuracy. Similarly, sparse representation methods have achieved 95.4% accuracy in classifying diffuse lung disease patterns on HRCT images, effectively handling the high variety and complexity of patterns[22].

2.1.3 Explainable AI in Medical Imaging

Clinical radiology requires an understandable method for determining how the disease occurs. The black-box nature of deep neural networks creates significant barriers for clinical adaptation. Physicians need to understand diagnostic reasoning to ensure accountability and patient safety. [23]. To address these challenges, a range of explainable artificial intelligence (XAI) techniques have been developed to provide insight into model decision-making. Predominant XAI approaches in medical imaging include[24]:

Perturbation-Based Methods: These techniques evaluate feature importance by systematically altering specific image regions and observing how these changes affect the model’s output.

Gradient-Based Methods: These methods highlight image areas that most significantly utilizing backpropagation gradients influence a classification. Grad-CAM is particularly predominant because it is applicable to nearly any network architecture [25].

Decomposition-Based Methods: Techniques like Layer-Wise Relevance Propagation (LRP) distribute the output of the final layer back through the network recursively. This process calculates a relevance score for each neuron, generating a heatmap that provides a clearer distinction of feature contribution.

These explainability techniques greatly increase clinician comprehension and confidence in adapting AI-driven diagnostic results, making it easier to integrate with conventional radiological workflows while maintaining regulatory compliance [25].

2.1.4 Volumetric Clinical Feature Extraction

The core of pulmonary diagnostic analysis lies in the identification and quantification of specific radiographic patterns. GGO is a hallmark sign in chest CT, characterized by hazy increased attenuation that does not obscure the underlying vascular markings. It is a non-specific finding that can represent fluid, fibrosis, or partial collapse of alveoli, and it is particularly prevalent in the early stages of diseases like COVID-19 or certain types of non-small cell lung cancer.

Volumetric analysis allows for the quantification of these features beyond simple 2D identification. By segmenting the lung parenchyma and using HU thresholding, the percentage of lung tissue affected can be calculated by specific patterns. A critical metric in this domain is the C/G ratio. Clinical data shows that as GGO progresses to consolidation, the severity of the patient’s condition increases[26].

2.2 Gaps in Literature

The proposed solution is designed to address several gaps in the current literature, as summarized in table 2.1:

Table 2.1: Gaps addressed by proposed approach vs existing solutions

Aspect	Existing Solutions	Proposed Solution
Interpretability	Post-hoc explanations (Grad-CAM, LIME, SHAP); Black-box features with no semantic meaning[24]	Inherently interpretable dictionary atoms representing concrete anatomical structures (nodules, GGO, consolidations)[16]
3D Volumetric Analysis	Limited integration of interpretability with full 3D data; mostly 2D slice-by-slice or post-hoc 3D explanations[27]	Full 3D volumetric analysis with interpretable sparse representations[22]
Feature Representation	Black-box features with no semantic meaning[12]	Dictionary atoms represent concrete, visualizable anatomical structures (nodules, GGO, consolidations)[19]
Label Consistency	Standard K-SVD focuses only on reconstruction error, lacks discriminative power[18]	Dictionary learning with sparse representations is more efficient than deep 3D CNNs[21]
Clinical Metrics Integration	Often limited to classification accuracy; lacks integration of domain-specific quantitative metrics[26]	Integrates HU-based segmentation for GGO/consolidation detection and C/G ratio calculation[28]
Clinical Metrics Integration	Deep learning requires large labeled datasets[13]	Dictionary learning is effective with limited labeled data; particularly valuable for rare diseases[16]
Handling Class Imbalance	Requires synthetic data generation, complex loss weighting, or meta-learning[29]	Sparse dictionary learning naturally handles few-shot scenarios for rare diseases[21]

2.2.1 Inherently Interpretable Architectures

While 3D CNNs have demonstrated superior performance for lung disease classification from CT scans, the majority of explainability research has focused on 2D imaging modalities (chest X-rays) or individual CT slices analyzed independently [27]. There remains a significant gap in developing inherently interpretable feature representations specifically designed for full 3D volumetric data [30].

Modern XAI methodologies (e.g., 3D Grad-CAM) frequently utilize post-hoc explanations to visualize important regions but do not provide semantically meaningful feature representations that clinicians can directly relate to pathological patterns [30]. These approaches explain model decisions after they are made but fail to tackle their underlying black-box nature. The extension of 3D saliency maps to interpretable concepts such as specific lesion types, textural patterns, or anatomical landmarks remains largely unexplored. This gap is particularly significant, because pulmonary pathology often manifests as spatial patterns spanning multiple CT slices (e.g., centrilobular versus panlobular emphysema distributions).

Sparse dictionary learning presents an alternative approach where learned dictionary elements can represent interpretable features such as nodules, GGO, consolidations, etc [16]. Every classification decision links to particular, tangible dictionary components, offering inherent rather than post-implementation clarity. Despite the integration of sparse dictionary learning principles with contemporary 3D deep learning models, significant exploration in 3D CT analysis is still lacking [31]. The existing hybrid methods, such as deep dictionary learning[32], have mainly been applied to 2D natural images rather than to volumetric medical imaging.

2.2.2 Frozen Dictionary Learning for Pathological Differentiation

A primary innovation of the proposed system is the utilization of frozen dictionary learning. In conventional dictionary learning, all atoms are updated simultaneously, which can lead to a phenomenon where normal atoms inadvertently absorb pathological features [18], reducing the model’s ability to clearly isolate anomalies. Frozen dictionary learning addresses this by employing a staged training strategy that partitions the dictionary into anatomically standard and pathologically specific components [33].

The first phase involves learning a base dictionary using only healthy lung data [34]. These atoms capture the expected appearance of lung parenchyma, bronchi, and pulmonary vasculature. Once this base dictionary is optimized, it is frozen (its atoms are held constant and are no longer updated) [33]. In the second phase, additional dictionary atoms are introduced and trained specifically on anomalous data such as CT scans containing viral pneumonia or adenocarcinoma. Because the standard anatomical features are already accounted for by the frozen atoms, the new atoms are forced to capture only the unique variations or lesions that represent the disease state.

When a query 3D patch is analyzed, the resulting sparse coefficient vector reveals exactly which atoms [35]. If the representation relies heavily on the anomalous portion of the dictionary, the clinician can identify the specific type of pathology present by mapping those atoms back to their visual counterparts in the original signal domain. This feature separation ensures that the diagnostic process is transparent and verifiable, moving away from the hidden logic of deep neural networks.

2.2.3 Handling Class Imbalance with Interpretable Models

Rare lung pathologies such as certain cancer subtypes (e.g., large cell carcinoma, carcinoid tumors), uncommon ILD patterns (e.g., pleuroparenchymal fibroelastosis), and emerging infectious diseases present severe class [13]. Deep learning research has addressed this through weighted loss functions, oversampling strategies, synthetic data generation (e.g., progressive growing GANs), and meta-learning approaches [29]. However, the application of class imbalance mitigation techniques within sparse representation frameworks for 3D CT analysis is underexplored [16]. Dictionary learning methods that can effectively learn representative atoms for rare categories while maintaining interpretability for common pathologies require further development [35]. The few-shot learning capabilities of sparse dictionary approaches offer potential advantages for rare disease applications, but have not been systematically evaluated in pulmonary imaging contexts.

Chapter 3

Methodology

The research consists of 7 main phases as shown in figure 3.1:




Figure 3.1: Overall Methodology

3.1 Data acquisition and pre-processing

We selected two publicly available 3D chest CT datasets, CT-RATE[36] and ReXGroundingCT[37], for training and evaluation due to their complementary nature. Their characteristics are summarized in table 3.1:

Table 3.1: Characteristics of CT-RATE and RexGroundingCT datasets

Feature	CT-RATE dataset	RexGroundingCT dataset
Type of Data	3D non-contrast chest CT volumes with corresponding radiology text reports and multi-abnormality labels	Subset of CT-RATE scans paired with free-text findings and pixel-level 3D segmentation masks
Number of Scans	25,692 non-contrast chest CT studies (50,188 volumes after reconstruction)	3,142 non-contrast chest CT scans
Number of patients	21,304 unique patients.	Subset selected from CT-RATE
Associated Reports	Full radiology reports (findings and impressions) with structured multi-abnormality labels (18 categories)	Standardized radiology reports from CT-RATE used to extract free-text findings
Annotations	Text-image pairs with structured labels; no voxel-wise ground truth segmentations	Manual pixel-level 3D segmentations of findings extracted from reports (8,028 findings; 16,301 entities)

Data pre processing corrections such as intensity normalization and spacing adjustments have already been implemented directly in the dataset. If further pre-processing is required, standard techniques such as resampling to a uniform voxel size, intensity normalization, and lung segmentation, shall be implemented before feeding the data into the model. Both datasets are released under MIT licenses that allow use for non-commercial academic research.

3.2 Model Training

Model training will be implemented using a two-phase frozen dictionary learning approach[38] to capture both normal and pathological patterns in chest CT scans. The dataset will be divided into train, test, and validation sets. The approach first learns a dictionary from healthy lung scans to model normal anatomical patterns, and then, while keeping this normal representation fixed, incrementally learns additional atoms to capture disease specific pathological features during sparse coding.

We extend the dictionary learning approach using LC-KSVD to improve the discriminative power of the learned sparse representations. The label consistency constraint will produce similar sparse codes for CT samples belonging to the same disease category. LC-KSVD will be formulated to jointly optimize three key objectives:

1. Accurate reconstruction of CT image patches using the learned dictionary.
2. Consistency of sparse representations within the same disease class.
3. Effective classification of lung conditions based on the sparse codes.

Appropriate weighting will be applied to balance these objectives and hyperparameters will be selected through validation experiments.

3.3 Classification and Disease Diagnosis




Figure 3.2: Model Training Workflow

Disease classification will be performed using a low-order linear classifier operating on the sparse representations produced by the learned dictionary. In particular, classification is achieved by applying a linear projection to the sparse codes, followed by a softmax-based decision rule. This approach enables multi-class prediction while maintaining a small number of trainable parameters. Class imbalance is addressed through the use of class-weighted loss functions during training. The classification module will be implemented in Python using scikit-learn, and its performance will be evaluated across all disease categories defined in the CT-RATE dataset.

3.4 Interpretability Analysis and Visualization

To improve interpretability and clinical relevance, the learned dictionary atoms will be visualized and analyzed through an interactive user interface utilizing 3D rendering techniques. Each dictionary atom will be mapped back to the original CT image space, allowing us to see which anatomical regions or pathological patterns the model has learned to represent. This mapping is guided by the ReXGroundingCT annotations, which provide precise spatial localization of abnormalities within the 3D imaging volume. Interpretability is measured using semantic coherence, atom-to-pattern correspondence, reconstruction fidelity, and segmentation accuracy, which together evaluate how meaningfully each learned atom represents clinically relevant structures when mapped back to the original 3D CT space.

3.5 Volumetric Feature Extraction and Quantification

Volumetric quantification of pulmonary patterns will be performed using HU-based segmentation to differentiate between GGO and consolidation regions. GGO will be identified using HU thresholds in the range of -700 to -250 HU[28], where increased attenuation preserves underlying vascular structures. Consolidation will be segmented using threshold greater than -250 HU[28], corresponding to homogeneous opacification that obscures bronchial and vascular margins. The segmented regions will be extracted from the 3D CT volumes to calculate the total volume of each pattern type.

The C/G ratio will be computed by dividing the total volume of consolidation by the total volume of GGO across the entire lung parenchyma. This quantitative metric provides an objective measure of disease severity and progression, with higher ratios indicating more advanced pathological states. The calculated C/G ratios will be compared against model predictions and clinical thresholds to validate the diagnostic utility of the learned dictionary representations. These volumetric features will complement the sparse coding analysis, enabling a more comprehensive evaluation of the model’s ability to capture clinically relevant patterns.

3.6 Baseline Comparisons and Benchmarking

The proposed method will be compared against well established baseline approaches commonly used in medical imaging research. These baselines are chosen to represent both traditional machine learning techniques and modern deep learning models. Standard K-SVD is included as a baseline to highlight the advantage of incorporating label information through LC-KSVD. By comparing these two sparse representation methods, the contribution of label consistent dictionary learning to disease classification performance can be assessed.

In addition, several 3D CNN architectures will be used as deep learning baselines, as they are considered state-of-the-art for volumetric CT analysis[7]. These

models are capable of learning complex spatial patterns across multiple slices and therefore provide a strong benchmark against which the proposed method can be evaluated. Existing implementations of these networks will be adopted to ensure reliability and consistency. All models will be trained and tested using the same dataset, although training strategies may differ due to the higher data and computational demands of CNN-based approaches.

All methods will be implemented using suitable and widely adopted frameworks, including scikit-learn for classical machine learning and sparse coding methods, and PyTorch or TensorFlow for deep learning based models.

3.7 Evaluation Metrics and Validation Strategy

The performance of the proposed model will be evaluated using accuracy, precision, recall, F1-score, and AUC-ROC, which are standard and widely accepted metrics in medical image analysis. Greater emphasis will be placed on accuracy and recall, as they are especially important in a clinical context to ensure overall reliability and to minimize missed disease cases, while the remaining metrics provide additional insight into class-wise performance and discriminative ability.

Chapter 4

Timeline and Resources Required

4.1 Timeline

The project will be carried out within 2 semesters (24 weeks) according to the timeline given in table [4.1](#).

4.2 Resources Required

The project requires several Hardware and Software resources as summarized in table [4.2](#).

Table 4.1: Overall Project Time Plan

Project Task	Semester 07												Semester 08											
	1	2	3	4	5	6	7	8	9	10	11	12	1	2	3	4	5	6	7	8	9	10	11	12
Dataset acquisition & preprocessing																								
Implement frozen dictionary + LC-KSVD & initial experiments																								
Implement volumetric analysis & feature extraction																								
Clinical validation & comparative experiments																								
Final experiments & visualization																								

Table 4.2: Resources Required for the Project

Resource Category	Specification	Availability	Cost
Hardware			
GPU	NVIDIA GPU (minimum 16GB VRAM) for training 3D CNN baselines and processing volumetric CT data	Institutional	N/A
CPU	Multi-core processor (minimum 8 cores) for dictionary learning algorithms	Institutional	N/A
RAM	Minimum 32GB for handling large 3D CT volumes in memory	Institutional	N/A
Storage	Minimum 500GB for storing CT-RATE and ReXGroundingCT datasets	Existing	N/A
Software & Libraries			
ML Frameworks	PyTorch or TensorFlow for 3D CNNs scikit-learn for SVM classification	Open Source	Free
Medical Imaging	SimpleITK or NiBabel for CT processing PyDicom for DICOM handling scikit-image for preprocessing	Open Source	Free
Visualization	Three.js or VTK for 3D visualization ITK-SNAP or 3D Slicer for qualitative analysis	Open Source	Free
Datasets			
CT-RATE	50,188 3D chest CT volumes with radiology reports	Open Source (MIT License)	Free
ReXGroundingCT	3,142 chest CT scans with annotated 3D segmentation masks	Open Source (MIT License)	Free

Chapter 5

Conclusion

This proposal presents a sparse dictionary learning framework for interpretable lung disease diagnosis from 3D chest CT scans. By employing frozen dictionary learning with LC-KSVD, the proposed approach addresses fundamental limitations of existing deep learning methods: computational overhead, data scarcity, and lack of clinical interpretability. The two-phase training strategy ensures clear separation between normal anatomical features and pathological patterns, enabling transparent diagnostic reasoning that clinicians can verify and trust.

The integration of volumetric clinical metrics, specifically HU-based segmentation and C/G ratio calculation, bridges the gap between machine learning predictions and established radiological practice. Dictionary atoms provide semantically meaningful representations that map directly to observable pathological structures, moving beyond post-hoc explanations toward inherently interpretable models.

Benchmarking against K-SVD and 3D CNN baselines will validate the effectiveness of label-consistent dictionary learning for medical image classification. The use of publicly available datasets (CT-RATE and ReXGroundingCT) ensures reproducibility and facilitates future research extensions. Ultimately, this work contributes to the development of open book AI systems that prioritize clinical transparency alongside diagnostic performance, offering a viable alternative to black-box deep learning approaches in high-stakes medical applications.

References

- [1] Y. Zhai, C. Zhu, T. Zhu, W. Song, Y. Tang, L. Jiang, F. Ruan, Z. Xu, L. Li, X. Fu, D. Liu, A. Chen, and Q. Wu, “Global, regional, and national burden of chronic respiratory diseases, 1990–2021 and predictions to 2035: analysis of data from the global burden of disease study 2021,” *Annals of Medicine*, vol. 57, no. 1, p. 2530225, 2025, PMID: 40627457; PMCID: PMC12239243.
- [2] M. O. Wielpütz, C. P. Heußel, F. J. F. Herth, and H.-U. Kauczor, “Radiological diagnosis in lung disease: Factoring treatment options into the choice of diagnostic modality,” *Deutsches Ärzteblatt International*, vol. 111, no. 11, pp. 181–187, Mar 2014.
- [3] H. Kasban, M. A. M. El-Bendary, and D. H. Salama, “A comparative study of medical imaging techniques,” *International Journal of Information Science and Intelligent System*, vol. 4, no. 2, pp. 37–58, 2015.
- [4] D. S. Gierada, D. G. Politte, J. Zheng, K. B. Schechtman, B. R. Whiting, K. E. Smith, T. Crabtree, D. Kreisel, A. S. Krupnick, G. A. Patterson, V. Puri, and B. F. Meyers, “Quantitative ct classification of lung nodules: Initial comparison of 2d and 3d analysis,” *Journal of Computer Assisted Tomography*, vol. 40, no. 4, pp. 589–595, 2016.
- [5] P. Maken and A. Gupta, “2d-to-3d: A review for computational 3d image reconstruction from x-ray images,” *Archives of Computational Methods in Engineering*, vol. 30, pp. 1–30, 07 2022.
- [6] R. Najjar, “Redefining radiology: A review of artificial intelligence integration in medical imaging,” *Diagnostics*, vol. 13, no. 17, 2023. [Online]. Available: <https://www.mdpi.com/2075-4418/13/17/2760>
- [7] S. Tiwari, G. Jain, D. K. Shetty, M. Sudhi, J. M. Balakrishnan, and S. R. Bhatta, “A comprehensive review on the application of 3d convolutional neural networks in medical imaging,” *Engineering Proceedings*, vol. 59, no. 1, p. 3, 2023. [Online]. Available: <https://www.mdpi.com/2673-4591/59/1/3>
- [8] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, “Grad-cam: Visual explanations from deep networks via gradient-based localization,” *International Journal of Computer Vision*, vol. 128, no. 2, pp. 336–359, Oct. 2019. [Online]. Available: <http://dx.doi.org/10.1007/s11263-019-01228-7>

- [9] M. Al Khalidah, M. Alqahtani *et al.*, “Deep learning models for CT image classification: a comprehensive literature review,” *Diagnostics*, vol. 14, no. 18, p. 2058, 2024.
- [10] N. Rai *et al.*, “An explainable AI approach for lung cancer detection using convolutional neural networks,” *arXiv preprint arXiv:2508.10196*, 2025.
- [11] S. P. Singh, L. Wang, S. Gupta, H. Goli, P. Padmanabhan, and B. Gulyás, “3d deep learning on medical images: A review,” 2020. [Online]. Available: <https://arxiv.org/abs/2004.00218>
- [12] G. Litjens, T. Kooi, B. E. Bejnordi, A. S’eppe, R. Clément, F. Ciompi, M. Ghafoorian, J. A. van der Laak, B. van Ginneken, and C. I. S’anchez, “A survey on deep learning in medical image analysis,” *Medical image analysis*, vol. 42, pp. 60–88, 2017.
- [13] Z. Liu, S. Wang, D. Dong, J. Wei, C. Fang, X. Zhou, K. Sun, D. Ni, Z. Liu, and Y. Jiang, “Deep learning for lung cancer diagnosis, prognosis, and immunotherapy: current progress and perspectives,” *Frontiers in Oncology*, vol. 12, 2022.
- [14] X. Li and et al., “Multi-label classification of lung diseases with densely connected pyramid deconvolutional network,” *Computers in Biology and Medicine*, 2023.
- [15] Y. Wang and et al., “Resnest50-based model for lung pathology classification with lung extraction,” *Biomedical Signal Processing and Control*, vol. 89, 2024.
- [16] R. Zhao *et al.*, “A survey on dictionary learning in medical image analysis,” *Medical Image Analysis*, 2014.
- [17] T. Ogawa and M. Haseyama, “Image inpainting based on sparse representations with a perceptual metric,” *EURASIP Journal on Advances in Signal Processing*, vol. 2013, p. 179, 12 2013.
- [18] M. Aharon, M. Elad, and A. Bruckstein, “K-svd: An algorithm for designing overcomplete dictionaries for sparse representation,” *IEEE Transactions on signal processing*, vol. 54, no. 11, pp. 4311–4322, 2006.
- [19] T. H. Vu, H. S. Mousavi, V. Monga, G. Rao, and U. K. A. Rao, “Histopathological image classification using discriminative feature-oriented dictionary learning,” *IEEE Transactions on Medical Imaging*, vol. 35, no. 3, pp. 738–751, 2016.
- [20] Q. Li, X. Wu, L. Xu, K. Chen, L. Yao, and Alzheimer’s Disease Neuroimaging Initiative, “Classification of alzheimer’s disease, mild cognitive impairment, and cognitively unimpaired individuals using multi-feature kernel discriminant dictionary learning,” *Frontiers in Computational Neuroscience*, vol. 11, p. 117, 2017. [Online]. Available: <https://www.frontiersin.org/journals/computational-neuroscience/articles/10.3389/fncom.2017.00117/full>

- [21] K. Liyanage, F. Ramezani, and B. Whitaker, *Covid-19 Chest CT Scan Image Classification Using LCKSVD and Frozen Sparse Coding*. Springer Nature Singapore, 01 2022, pp. 272–281.
- [22] W. Zhao, R. Xu, Y. Hirano, R. Tachibana, and S. Kido, “Classification of diffuse lung diseases patterns by a sparse representation based method on hrct images,” in *2013 35th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, 2013, pp. 5457–5460.
- [23] B. H. van der Velden, K. G. van Leeuwen, R. H. van der Sluijs, J. Gilbers, M. A. Ragusi, M. W. van den Brekel *et al.*, “Unveiling the black box: A systematic review of explainable artificial intelligence in medical image analysis,” *Computational and Structural Biotechnology Journal*, vol. 23, pp. 2642–2656, 2024.
- [24] D. Bhati, F. Neha, and M. Amiruzzaman, “A survey on explainable artificial intelligence (XAI) techniques for visualizing deep learning models in medical imaging,” *Journal of Imaging*, vol. 10, no. 10, p. 239, 2024.
- [25] G. Singh, A. Al-Huty *et al.*, “Explainable artificial intelligence for medical imaging systems using deep learning: a comprehensive review,” *Cluster Computing*, 2025.
- [26] W. R. Webb, N. L. Muller, and D. P. Naidich, *Fundamentals of High-Resolution Lung CT: Common Findings, Common Patterns, Common Diseases, and Differential Diagnosis*, 4th ed. Elsevier, 2014.
- [27] H. Li *et al.*, “Volumetric xai for 3d medical image analysis: A survey,” in *MICCAI Workshop*, 2021.
- [28] H. Tanaka, N. Tanabe, S. Chubachi, T. Maetani, Y. Shiraishi, H. Namkoong, T. Asakura, T. Shimada, S. Azekawa, S. Otake, K. Nakagawara, T. Fukushima, M. Watase, H. Terai, M. Sasaki, S. Ueda, Y. Kato, N. Harada, S. Suzuki, S. Yoshida, H. Tateno, Y. Yamada, M. Jinzaki, T. Hirai, Y. Okada, R. Koike, M. Ishii, N. Hasegawa, A. Kimura, S. Imoto, S. Miyano, S. Ogawa, T. Kanai *et al.*, “Consolidation-to-ground-glass opacity ratio on chest ct as a prognostic marker for critical outcomes in covid-19: a retrospective cohort study,” *BMC Pulmonary Medicine*, vol. 25, p. 550, 2025. [Online]. Available: <https://link.springer.com/article/10.1186/s12890-025-04030-z>
- [29] J. M. Johnson and T. M. Khoshgoftaar, “Survey on deep learning with class imbalance,” *Journal of Big Data*, vol. 6, no. 1, pp. 1–82, 2019.
- [30] S. Huang *et al.*, “A survey on explainable ai for 3d medical imaging,” *Medical Image Analysis*, 2022.
- [31] S. Gu, L. Zhang, W. Zuo, and X. Zhang, “Deep dictionary learning for sparse representation,” *CVPR*, 2018.
- [32] S. Tariyal, A. Majumdar, R. Singh, and M. Vatsa, “Deep dictionary learning,” *IEEE Access*, vol. 4, pp. 10 096–10 109, 2016.

- [33] V. Nika, C. Davatzikos *et al.*, “Change detection of medical images using dictionary learning techniques,” *Journal of Medical Imaging*, 2014.
- [34] Q. Xu, H. Yu, X. Mou, L. Zhang, J. Hsieh, and M. Kalra, “Low-dose x-ray ct reconstruction via dictionary learning,” *IEEE Transactions on Medical Imaging*, vol. 31, no. 4, pp. 1268–1277, 2012.
- [35] Z. Jiang, Z. Lin, and L. S. Davis, “Label consistent k-svd: Learning a discriminative dictionary for recognition,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 35, no. 11, pp. 2651–2664, 2013.
- [36] I. E. Hamamci, S. Er, C. Wang, F. Almas, A. G. Simsek, S. N. Esirgun, I. Dogan, O. F. Durugol, B. Hou, S. Shit, W. Dai, M. Xu, H. Reynaud, M. F. Dasdelen, B. Wittmann, T. Amiranashvili, E. Simsar, M. Simsar, E. B. Erdemir, A. Alanbay, A. Sekuboyina, B. Lafci, A. Kaplan, Z. Lu, M. Polacin, B. Kainz, C. Bluethgen, K. Batmanghelich, M. K. Ozdemir, and B. Menze, “Developing generalist foundation models from a multimodal dataset for 3d computed tomography,” 2025. [Online]. Available: <https://arxiv.org/abs/2403.17834>
- [37] M. Baharoon, L. Luo, M. Moritz, A. Kumar, S. E. Kim, X. Zhang, M. Zhu, M. H. Alabbad, M. S. Alhazmi, N. P. Mistry, L. Bijnens, K. R. Kleinschmidt, B. Chrisler, S. Suryadevara, S. S. D. Jaliparthi, N. M. Prudlo, M. D. Marino, J. Palacio, R. Akula, D. Zhou, H.-Y. Zhou, I. E. Hamamci, S. J. Adams, H. R. AlOmaish, and P. Rajpurkar, “Rexgroundingct: A 3d chest ct dataset for segmentation of findings from free-text reports,” 2025. [Online]. Available: <https://arxiv.org/abs/2507.22030>
- [38] M. Aharon, M. Elad, and A. Bruckstein, “K-svd: An algorithm for designing overcomplete dictionaries for sparse representation,” *IEEE Transactions on Signal Processing*, 2006.